

Project 1 - Regression Problem

Dataset Description

Link to the dataset used:

<https://www.kaggle.com/sameersmahajan/seattle-house-sales-prices>
(<https://www.kaggle.com/sameersmahajan/seattle-house-sales-prices>)

ID - Unique ID for each home sold

Date - Date of the home sale

Price - Price of the home sale

Bedrooms - Number of bedrooms

Bathrooms - Number of bathrooms (0.5 accounts for a room with a toilet but no shower)

Sqft_living - Square footage of the apartments

Sqft_lot - Square footage of the land space

Floors - Number of floors

Waterfront - A dummy variable for whether the apartment was overlooking the waterfront or not

View - Index from 0 to 4 of how good the view of the property was

Condition - An index from 1 to 5 on the condition of the apartment

Grade - An index from 1 to 13, where 1-3 falls short of building construction and design, 7 has an average level of construction and design, and 11-13 have a high quality level of construction and design

Sqft_above - The square footage of the interior housing space that is above ground level

Sqft_basement - The square footage of the interior housing space that is below ground level

Yr_built - The year the house was initially built

Yr_renovated - The year of the house's last renovation

Zipcode - What zipcode area the house is in

Lat - Latitude

Long - Longitude

Sqft_living15 - The square footage of interior housing living space for the nearest 15 neighbors

Sqft_lot15 - The square footage of the land lots of the nearest 15 neighbors

To print out the output of all codes in a cell

```
In [1]: from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
```

Loading Libraries

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.preprocessing import PolynomialFeatures
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.svm import SVR
from sklearn.svm import LinearSVR
from sklearn.model_selection import GridSearchCV
```

Loading Dataset

```
In [4]: bm = pd.read_csv("house_sales.csv")
```

Overview of Dataset

```
In [5]: bm.shape  
bm.head(10)
```

```
Out[5]: (21613, 21)
```

```
Out[5]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	water
0	7129300520	20141013T000000	221900	3	1.00	1180	5650	1.0	
1	6414100192	20141209T000000	538000	3	2.25	2570	7242	2.0	
2	5631500400	20150225T000000	180000	2	1.00	770	10000	1.0	
3	2487200875	20141209T000000	604000	4	3.00	1960	5000	1.0	
4	1954400510	20150218T000000	510000	3	2.00	1680	8080	1.0	
5	7237550310	20140512T000000	1225000	4	4.50	5420	101930	1.0	
6	1321400060	20140627T000000	257500	3	2.25	1715	6819	2.0	
7	2008000270	20150115T000000	291850	3	1.50	1060	9711	1.0	
8	2414600126	20150415T000000	229500	3	1.00	1780	7470	1.0	
9	3793500160	20150312T000000	323000	3	2.50	1890	6560	2.0	

10 rows × 21 columns



```
In [6]: bm.columns  
bm.dtypes
```

```
Out[6]: Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',  
              'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'grade',  
              'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode',  
              'lat', 'long', 'sqft_living15', 'sqft_lot15'],  
             dtype='object')
```

```
Out[6]: id                int64  
date                  object  
price                int64  
bedrooms             int64  
bathrooms            float64  
sqft_living          int64  
sqft_lot             int64  
floors              float64  
waterfront           int64  
view                 int64  
condition            int64  
grade               int64  
sqft_above           int64  
sqft_basement        int64  
yr_built             int64  
yr_renovated         int64  
zipcode             int64  
lat                 float64  
long                float64  
sqft_living15        int64  
sqft_lot15           int64  
dtype: object
```

In [7]: `bm.head().T`

Out[7]:

	0	1	2	3	
id	7129300520	6414100192	5631500400	2487200875	19544
date	20141013T000000	20141209T000000	20150225T000000	20141209T000000	20150218T000000
price	221900	538000	180000	604000	500000
bedrooms	3	3	2	4	4
bathrooms	1	2.25	1	3	3
sqft_living	1180	2570	770	1960	1960
sqft_lot	5650	7242	10000	5000	5000
floors	1	2	1	1	1
waterfront	0	0	0	0	0
view	0	0	0	0	0
condition	3	3	3	5	5
grade	7	7	6	7	7
sqft_above	1180	2170	770	1050	1050
sqft_basement	0	400	0	910	910
yr_built	1955	1951	1933	1965	1965
yr_renovated	0	1991	0	0	0
zipcode	98178	98125	98028	98136	98136
lat	47.5112	47.721	47.7379	47.5208	47.5208
long	-122.257	-122.319	-122.233	-122.393	-122.393
sqft_living15	1340	1690	2720	1360	1360
sqft_lot15	5650	7639	8062	5000	5000

Data Preprocessing

Converting the data types of columns accordingly for exploring the dataset easier

Columns


'id', 'bedrooms', 'bathrooms', 'floors', 'waterfront', 'view', 'condition', 'grade', 'yr_built', 'yr_renovated', 'zipcode' are categorical

In [8]: `bm[['id', 'bedrooms', 'bathrooms', 'floors', 'waterfront', 'view', 'condition', 'grade', 'yr_built', 'yr_renovated', 'zipcode']]`

Dropping the column 'id' as it does not affect the prediction of house prices.

```
In [9]: sns.scatterplot(x='id',y='price',data=bm)
```

```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9790e320>
```



```
In [10]: bm_1 = bm.drop(['id'],axis=1)
```

Creating and Handling Null Values

Null values are created randomly across the dataset

Creating a dataframe without column 'Price' so that null values are not created under the dependent column 'Price'

```
In [11]: bm_price = bm_1['price']  
bm_x = bm_1.drop(['price'],axis=1)  
bm_x.shape
```

```
Out[11]: (21613, 19)
```

```
In [12]: np.random.seed(0)  
bm_x = bm_x.mask(np.random.random(bm_x.shape) < .051)
```

```
In [13]: bm_x.isnull().sum()
```

```
Out[13]: date           1128  
bedrooms           1099  
bathrooms          1140  
sqft_living        1120  
sqft_lot           1069  
floors             1061  
waterfront         1096  
view               1119  
condition          1078  
grade              1087  
sqft_above         1127  
sqft_basement      1080  
yr_built           1031  
yr_renovated       1177  
zipcode            1128  
lat                1087  
long               1137  
sqft_living15      1097  
sqft_lot15         1123  
dtype: int64
```

```
In [14]: bm_x.isnull().sum().sum()
```

```
Out[14]: 20984
```

20984 null values (approx. 5-6% of whole dataset) are created across the dataset

```
In [15]: bm_x[bm_x.isnull().any(axis=1)]
```

```
Out[15]:
```

	date	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	cond
0	20141013T000000	3	1	1180.0	5650.0	1	0	0	
1	20141209T000000	3	2.25	2570.0	7242.0	2	0	0	
3	20141209T000000	4	3	1960.0	5000.0	1	0	0	
5	20140512T000000	4	NaN	5420.0	NaN	1	0	0	
6	20140627T000000	3	2.25	1715.0	6819.0	2	0	0	
8	20150415T000000	3	1	1780.0	7470.0	1	0	0	
9	20150312T000000	3	NaN	1890.0	6560.0	2	0	0	
10	20150403T000000	3	2.5	3560.0	9796.0	1	0	0	
11	20140527T000000	2	1	1160.0	6000.0	NaN	0	0	
12	20140528T000000	3	1	1430.0	19901.0	NaN	0	0	
13	20141007T000000	3	1.75	1370.0	9680.0	1	0	0	
14	20150312T000000	5	2	1810.0	4850.0	1.5	0	0	
16	20140731T000000	3	2	1890.0	14040.0	NaN	0	0	
18	20141205T000000	NaN	1	1200.0	9850.0	1	0	0	
20	20140514T000000	4	1.75	1620.0	4980.0	1	0	0	
21	20140826T000000	3	2.75	3050.0	44867.0	1	0	4	
23	20140516T000000	2	1.5	1070.0	9643.0	1	NaN	0	
26	20140626T000000	3	1.75	2450.0	2691.0	2	0	0	
27	20141201T000000	3	1	1400.0	1581.0	1.5	0	0	
28	20140624T000000	3	1.75	1520.0	6380.0	1	0	0	
30	20141110T000000	NaN	2.5	2320.0	3980.0	2	0	0	
35	20140613T000000	3	2.5	2300.0	3060.0	1.5	0	0	
36	20140528T000000	4	1	1660.0	34848.0	1	0	NaN	
38	NaN	4	1	1220.0	8075.0	1	0	0	
39	20140620T000000	NaN	2.5	2620.0	7553.0	2	0	0	
42	20140707T000000	5	NaN	3595.0	5639.0	2	0	0	
43	20141028T000000	3	1	1570.0	NaN	2	0	0	
44	20140729T000000	NaN	1	1280.0	9656.0	1	0	0	
48	20150428T000000	3	1.75	1250.0	5963.0	1	0	0	
49	NaN	3	2.5	2753.0	65005.0	1	1	2	
...	
21568	20150130T000000	4	3.5	3830.0	8963.0	2	0	NaN	
21569	NaN	2	2.5	980.0	1020.0	3	0	0	
21570	20140513T000000	3	2.5	1450.0	5008.0	1	0	0	
21572	20140707T000000	3	1.75	1140.0	1201.0	2	0	0	

	date	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	cond
21573	20140520T000000	4	3.5	3070.0	4684.0	2	0	0	
21574	20140507T000000	3	3	1680.0	1570.0	3	0	0	
21575	NaN	3	2.5	3087.0	5002.0	2	0	0	
21577	NaN	4	3.25	1900.0	2631.0	2	0	0	
21580	20141003T000000	3	3	2780.0	6000.0	2	0	NaN	
21581	20150504T000000	3	3	NaN	6000.0	2	0	0	
21583	20140610T000000	2	NaN	710.0	1157.0	NaN	0	0	
21585	20140828T000000	3	2.5	NaN	5000.0	2	0	0	
21586	NaN	2	2.5	1430.0	1201.0	3	0	0	
21587	20150305T000000	3	2.5	1520.0	1488.0	3	0	0	
21589	20140910T000000	3	NaN	2540.0	4760.0	2	0	0	
21590	20140514T000000	4	3.5	4910.0	9444.0	1.5	0	0	
21591	20141002T000000	4	2.75	2770.0	3852.0	NaN	0	0	
21593	20150317T000000	5	3.75	4170.0	8142.0	2	0	NaN	
21594	20141017T000000	4	2.75	2500.0	5995.0	NaN	0	0	
21597	20150421T000000	4	3.25	3410.0	10125.0	2	0	0	
21598	20141013T000000	4	2.5	3118.0	7866.0	2	0	2	
21600	20141015T000000	NaN	3.75	NaN	NaN	2	0	0	
21601	20150407T000000	3	2.5	1425.0	1179.0	3	0	0	
21605	20141014T000000	NaN	2.5	2520.0	6023.0	2	0	0	
21606	20150326T000000	4	3.5	3510.0	7200.0	2	0	0	
21608	20140521T000000	3	2.5	1530.0	1131.0	NaN	0	0	
21609	20150223T000000	4	2.5	NaN	5813.0	2	0	0	
21610	20140623T000000	2	0.75	1020.0	1350.0	NaN	0	0	
21611	20150116T000000	3	2.5	1600.0	2388.0	2	0	NaN	
21612	20141015T000000	2	0.75	1020.0	NaN	NaN	0	0	

13637 rows × 19 columns

There are 13,637 rows with null values. Hence, it is not advisable to drop all the rows with null values. Better method would be impute the null values wherever possible.

Combining the dataset with null values with the column "price"

```
In [16]: bm_final = pd.concat([bm_x,bm_price],axis=1)
```

```
In [17]: bm_final.shape  
bm_final.isnull().sum()
```

```
Out[17]: (21613, 20)
```

```
Out[17]: date          1128  
bedrooms          1099  
bathrooms          1140  
sqft_living        1120  
sqft_lot           1069  
floors             1061  
waterfront         1096  
view               1119  
condition          1078  
grade              1087  
sqft_above         1127  
sqft_basement      1080  
yr_built           1031  
yr_renovated       1177  
zipcode            1128  
lat                1087  
long               1137  
sqft_living15      1097  
sqft_lot15         1123  
price              0  
dtype: int64
```

Creating two columns from the 'date' column - sale year and month

```
In [18]: bm_final['sale_year'] = bm_final['date'].str[:4]
```

```
In [19]: bm_final['sale_month'] = bm_final['date'].str[4:6]
```

```
In [20]: bm_final.sale_year.fillna('nan',inplace=True)  
bm_final.sale_month.fillna('nan',inplace=True)
```

```
In [21]: bm_final['sale_year'].value_counts()
bm_final['sale_year'].unique()
bm_final['sale_month'].value_counts()
bm_final['sale_month'].unique()
```

```
Out[21]: 2014    13890
         2015     6595
         nan     1128
         Name: sale_year, dtype: int64
```

```
Out[21]: array(['2014', '2015', 'nan'], dtype=object)
```


```
Out[21]: 05    2282
         04    2110
         06    2083
         07    2081
         08    1851
         10    1779
         03    1769
         09    1682
         12    1404
         11    1330
         02    1188
         nan    1128
         01     926
         Name: sale_month, dtype: int64
```

```
Out[21]: array(['10', '12', '02', '05', '06', '01', '04', '03', '07', '08', '11',
                'nan', '09'], dtype=object)
```

'Sale_Year' and 'Sale_Month'


```
In [22]: sns.boxplot(x='sale_year',y='price',data=bm_final)
```

```
Out[22]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd99d7ad68>
```



```
In [23]: sns.boxplot(x='sale_month',y='price',data=bm_final)
```

```
Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9ac1a4a8>
```



```
In [24]: bm_month = bm_final[['sale_month','price']]
bm_month.groupby('sale_month').mean()
```

```
Out[24]:
```

	price
sale_month	
01	522241.234341
02	510867.186027
03	544219.409836
04	559923.530806
05	550857.780456
06	559200.218435
07	544410.164344
08	537774.264722
09	526874.860285
10	538100.658797
11	519752.912782
12	525533.121083
nan	545605.439716

We can infer that the mean price of the houses over the month does not vary a lot. Hence, dropping the null values in the column 'sale_month'

```
In [25]: bm_final = bm_final[bm_final.sale_month != 'nan']
```

```
In [26]: bm_year = bm_final[['sale_year', 'price']]
bm_year.groupby('sale_year').mean()
```

```
Out[26]:
```

	price
sale_year	
2014	539228.268035
2015	540955.486277

We can observe that the mean price of houses sold in 2015 is greater than those sold in 2014

From the analysis of "sale_year" column, we get to know that sales are made between the years 2014-2015. Also, from the plot, we can infer that the sales price are similar to those houses sold at 2015.

It would be advisable to impute the 'nan' values with 2015. But, before that we have to check the "year built" column so as to ensure year built is 'nan' or not for the rows with 'nan' under sale_year column

```
In [27]: bm_final[(bm_final['sale_year']=='nan') & (bm_final['yr_built']=='nan')]
```

```
Out[27]:
```

date	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	...
0 rows × 22 columns										

No rows with 'nan' in both columns 'yr_built' and 'sale_year'. Therefore, replacing null under 'sale_year' with '2015'

```
In [28]: bm_final['sale_year'] = bm_final['sale_year'].replace({'nan': '2015'})
```

Only 2 years are present: 2014, 2015

Hence, assigning 0 and 1 respectively

```
In [29]: bm_final['sale_year'] = bm_final['sale_year'].replace({'2014': '0'})
bm_final['sale_year'] = bm_final['sale_year'].replace({'2015': '1'})
```

```
In [30]: bm_final['sale_year'].value_counts()
```

```
Out[30]: 0    13890
1     6595
Name: sale_year, dtype: int64
```

'Bedrooms' and 'Bathrooms'

```
In [31]: bm_final.bedrooms.fillna('nan', inplace=True)
bm_final.bathrooms.fillna('nan', inplace=True)
```

```
In [32]: bm_final['bedrooms'].unique()
bm_final['bedrooms'].value_counts()
bm_final['bathrooms'].unique()
bm_final['bathrooms'].value_counts()
```

```
Out[32]: array([3, 2, 4, 5, 'nan', 1, 6, 7, 0, 8, 9, 11, 10, 33], dtype=object)
```

```
Out[32]: 3      8863
4      6199
2      2452
5      1431
nan     1046
6       243
1       180
7        36
8        13
0        12
9         5
10        3
33         1
11         1
Name: bedrooms, dtype: int64
```

```
Out[32]: array([1.0, 2.25, 3.0, 2.0, 'nan', 1.5, 2.5, 1.75, 2.75, 3.25, 4.0, 3.5,
0.75, 4.75, 5.0, 4.25, 4.5, 3.75, 0.0, 1.25, 5.25, 6.0, 0.5, 5.5,
6.75, 5.75, 8.0, 7.5, 7.75, 6.25, 6.5], dtype=object)
```

```
Out[32]: 2.5      4845
1.0      3466
1.75     2747
2.25     1837
2.0      1726
1.5      1286
nan       1073
2.75     1071
3.0        671
3.5        663
3.25       517
3.75       139
4.0        119
4.5         87
4.25        73
0.75        67
4.75        19
5.0         19
5.25        13
0.0          9
5.5          9
1.25         7
6.0          6
5.75         4
0.5          4
8.0          2
6.75         2
7.75         1
6.25         1
```

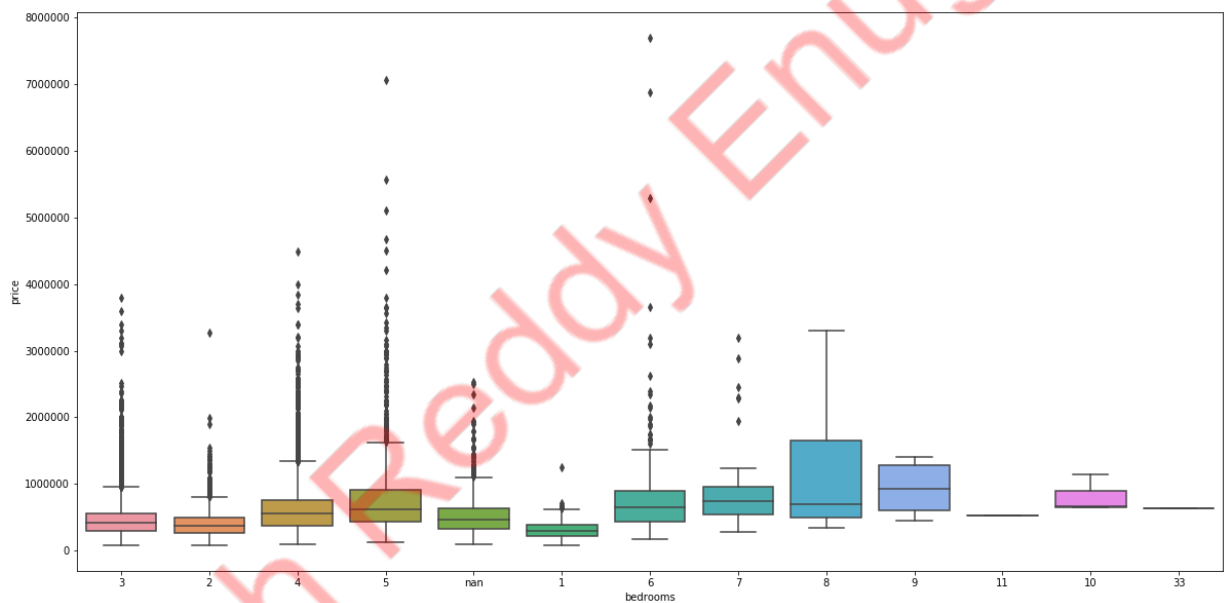
```
7.5      1
6.5      1
Name: bathrooms, dtype: int64
```

From above description of the number of bedrooms and bathrooms, we are able to find that there are houses with zero bedrooms and bathrooms, which is practically not possible. Hence, those rows can be dropped.

```
In [33]: bm_final = bm_final[bm_final.bedrooms != 0]
bm_final = bm_final[bm_final.bathrooms != 0]
```


```
In [34]: a4_dims = (20,10)
fig, ax = plt.subplots(figsize=a4_dims)
sns.boxplot(x='bedrooms',y='price',data=bm_final)
```

```
Out[34]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9afb3748>
```



```
In [35]: a4_dims = (20,10)
fig, ax = plt.subplots(figsize=a4_dims)
sns.boxplot(x='bathrooms',y='price',data=bm_final)
```

```
Out[35]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9ac40438>
```



Based on the plots, we can infer that the null values for the bedrooms can be replaced with '3' and for bathrooms with '2.25'

```
In [36]: bm_final['bedrooms'] = bm_final['bedrooms'].replace({'nan':3})
bm_final['bathrooms'] = bm_final['bathrooms'].replace({'nan':2.25})
```



```
In [37]: bm_final['bedrooms'].unique()
bm_final['bedrooms'].value_counts()
bm_final['bathrooms'].unique()
bm_final['bathrooms'].value_counts()
bm_final.shape
```

```
Out[37]: array([ 3,  2,  4,  5,  1,  6,  7,  8,  9, 11, 10, 33], dtype=int64)
```

```
Out[37]: 3      9908
4      6199
2      2452
5      1431
6       243
1       178
7        36
8        13
9         5
10        3
11        1
33        1
Name: bedrooms, dtype: int64
```

```
Out[37]: array([1. , 2.25, 3. , 2. , 1.5 , 2.5 , 1.75, 2.75, 3.25, 4. , 3.5 ,
0.75, 4.75, 5. , 4.25, 4.5 , 3.75, 1.25, 5.25, 6. , 0.5 , 5.5 ,
6.75, 5.75, 8. , 7.5 , 7.75, 6.25, 6.5 ])
```

```
In [38]: bm_final.isnull().sum()
```

```
Out[38]: date           0
bedrooms             0
bathrooms            0
sqft_living        1062
sqft_lot           1022
floors             1000
waterfront         1025
view               1078
condition          1019
grade              1031
sqft_above         1079
sqft_basement      1030
yr_built           971
yr_renovated       1102
zipcode            1060
lat                1027
long               1078
sqft_living15      1041
sqft_lot15         1059
.
```

"sqft_living"

```
In [39]: bm_final['sqft_living'].isnull().sum()  
bm_final['sqft_living'].describe()  
bm_final['sqft_living'].median()
```

Out[39]: 1062

```
Out[39]: count    19408.000000  
mean       2081.284367  
std        917.777180  
min         370.000000  
25%       1430.000000  
50%       1920.000000  
75%       2550.000000  
max      13540.000000  
Name: sqft_living, dtype: float64
```

Out[39]: 1920.0


```
In [40]: a4_dims = (20,10)  
fig, ax = plt.subplots(figsize=a4_dims)  
sns.scatterplot(x='sqft_living',y='price',data=bm_final)
```

Out[40]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9b4807f0>



```
In [41]: sns.boxplot(x='sqft_living',data=bm_final)
```

```
Out[41]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9ad55ac8>
```



Since, we are replacing the null values with median as there are outliers

```
In [42]: bm_final['sqft_living']=bm_final['sqft_living'].fillna(bm_final['sqft_living'].median())
```

```
In [43]: bm_final['sqft_living'].isnull().sum()
```

```
Out[43]: 0
```

```
In [44]: bm_final.isnull().sum()
```

```
Out[44]: date           0
bedrooms           0
bathrooms          0
sqft_living         0
sqft_lot          1022
floors             1000
waterfront         1025
view               1078
condition          1019
grade              1031
sqft_above         1079
sqft_basement      1030
yr_built           971
yr_renovated       1102
zipcode            1060
lat                1027
long               1078
sqft_living15      1041
sqft_lot15         1059
.
```

'sqft_lot'

```
In [45]: bm_final['sqft_lot'].isnull().sum()
bm_final['sqft_lot'].describe()
bm_final['sqft_lot'].median()
```


```
Out[45]: 1022
```

```
Out[45]: count      1.944800e+04
mean        1.506629e+04
std         4.177920e+04
min         5.200000e+02
25%         5.053750e+03
50%         7.620000e+03
75%         1.070400e+04
max         1.651359e+06
Name: sqft_lot, dtype: float64
```

```
Out[45]: 7620.0
```

```
In [46]: a4_dims = (20,10)
fig, ax = plt.subplots(figsize=a4_dims)
sns.scatterplot(x='sqft_lot',y='price',data=bm_final)
```

```
Out[46]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9ada3128>
```



Since, outliers are present, imputing the null values with median of the column 'sqft_lot'

```
In [47]: bm_final['sqft_lot']=bm_final['sqft_lot'].fillna(bm_final['sqft_lot'].median())
bm_final['sqft_lot'].isnull().sum()
```

```
Out[47]: 0
```

```
In [48]: bm_final.isnull().sum()
```

```
Out[48]: date                0
bedrooms                   0
bathrooms                  0
sqft_living                0
sqft_lot                  0
floors                     1000
waterfront                 1025
view                      1078
condition                 1019
grade                     1031
sqft_above                1079
sqft_basement             1030
yr_built                   971
yr_renovated              1102
zipcode                   1060
lat                       1027
long                     1078
sqft_living15             1041
sqft_lot15                1059
.
```

'floors'

```
In [49]: bm_final.floors.fillna('nan',inplace=True)
```

```
In [50]: bm_final['floors'].value_counts()
bm_price_floors = bm_final[['floors','price']]
bm_price_floors.groupby('floors').mean()
```

```
Out[50]: 1.0    9602
2.0    7436
1.5    1731
nan    1000
3.0     549
2.5     145
3.5        7
Name: floors, dtype: int64
```

```
Out[50]:
```

	price
floors	
1.0	4.415906e+05
1.5	5.584669e+05
2.0	6.483285e+05
2.5	1.069188e+06
3.0	5.889418e+05
3.5	9.102143e+05
nan	5.388322e+05

We can see from the above table that the mean of those 'nan' floors is close to that of floors with

1.5

Hence, it is advisable to impute the nan's with 1.5

```
In [51]: bm_final['floors'] = bm_final['floors'].replace({'nan':1.5})
```

```
In [52]: bm_final['floors'].value_counts()
```

```
Out[52]: 1.0    9602
         2.0    7436
         1.5    2731
         3.0     549
         2.5     145
         3.5        7
         Name: floors, dtype: int64
```

```
In [53]: bm_final.isnull().sum()
```

```
Out[53]: date            0
         bedrooms        0
         bathrooms       0
         sqft_living      0
         sqft_lot         0
         floors           0
         waterfront     1025
         view            1078
         condition       1019
         grade           1031
         sqft_above       1079
         sqft_basement    1030
         yr_built         971
         yr_renovated     1102
         zipcode          1060
         lat              1027
         long             1078
         sqft_living15    1041
         sqft_lot15       1059
         price            0
         sale_year        0
         sale_month       0
         dtype: int64
```


'waterfront'

```
In [54]: bm_final.waterfront.fillna('nan',inplace=True)
         bm_final['waterfront'].value_counts()
```

```
Out[54]: 0      19302
         nan    1025
         1      143
         Name: waterfront, dtype: int64
```

```
In [55]: sns.boxplot(x='waterfront',y='price',data=bm_final)
```

```
Out[55]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9ae1cdd8>
```



Since, for both 'nan' and '0', they have the similar price ranges. We are imputing the nan's with 0

```
In [56]: bm_final['waterfront'] = bm_final['waterfront'].replace({'nan':0})  
bm_final['waterfront'].value_counts()
```

```
Out[56]: 0    20327  
         1     143  
         Name: waterfront, dtype: int64
```



```
In [57]: bm_final.isnull().sum()
```

```
Out[57]: date                0
bedrooms                    0
bathrooms                   0
sqft_living                  0
sqft_lot                     0
floors                       0
waterfront                   0
view                        1078
condition                   1019
grade                       1031
sqft_above                   1079
sqft_basement                1030
yr_built                     971
yr_renovated                 1102
zipcode                     1060
lat                          1027
long                        1078
sqft_living15                1041
sqft_lot15                   1059
price                        0
sale_year                    0
sale_month                   0
dtype: int64
```

'view'


```
In [58]: bm_final.view.fillna('nan',inplace=True)
```

```
In [59]: bm_final['view'].value_counts()
```

```
Out[59]: 0      17520
nan      1078
2         846
3         444
1         299
4         283
Name: view, dtype: int64
```

```
In [60]: sns.boxplot(x='view',y='price',data=bm_final)
```

```
Out[60]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9b06ecc0>
```



Since, for both 'nan' and '0', they have the similar price ranges. We are imputing the nan's with 0

```
In [61]: bm_final['view'] = bm_final['view'].replace({'nan':0})
bm_final['view'].value_counts()
```

```
Out[61]: 0    18598
         2     846
         3     444
         1     299
         4     283
         Name: view, dtype: int64
```

```
In [62]: bm_final.isnull().sum()
```

```
Out[62]: date            0
bedrooms              0
bathrooms            0
sqft_living          0
sqft_lot             0
floors              0
waterfront          0
view                0
condition           1019
grade              1031
sqft_above          1079
sqft_basement       1030
yr_built            971
yr_renovated        1102
zipcode            1060
lat                1027
long              1078
sqft_living15       1041
sqft_lot15          1059
.
```


'condition'

```
In [63]: bm_final.condition.fillna('nan',inplace=True)
bm_final['condition'].value_counts()
```

```
Out[63]: 3      12584
4       5147
5       1544
nan      1019
2        147
1         29
Name: condition, dtype: int64
```

```
In [64]: sns.boxplot(x='condition',y='price',data=bm_final)
```

```
Out[64]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9b543208>
```



```
In [65]: bm_price_condition = bm_final[['condition','price']]
bm_price_condition.groupby('condition').mean()
```

```
Out[65]:
```

	price
condition	
1	341067.241379
2	318611.489796
3	540408.542832
4	522513.078104
5	611373.721503
nan	550373.646712

Since, for both 'nan' and '3', they have the similar price ranges. We are imputing the nan's with 3

```
In [66]: bm_final['condition'] = bm_final['condition'].replace({'nan':3})
bm_final['condition'].value_counts()
bm_final.isnull().sum()
```

```
Out[66]: 3    13603
4      5147
5      1544
2       147
1        29
Name: condition, dtype: int64
```

```
Out[66]: date          0
bedrooms             0
bathrooms            0
sqft_living          0
sqft_lot             0
floors              0
waterfront          0
view                0
condition            0
grade               1031
sqft_above           1079
sqft_basement        1030
yr_built             971
yr_renovated         1102
zipcode              1060
lat                  1027
long                 1078
sqft_living15        1041
sqft_lot15           1059
price                0
sale_year            0
sale_month           0
dtype: int64
```


'grade'

```
In [67]: bm_final.grade.fillna('nan',inplace=True)
bm_final['grade'].value_counts()
```

```
Out[67]: 7      8098
8      5477
9      2330
6      1818
nan     1031
10     1027
11      352
5       218
12       79
4        26
13       13
3         1
Name: grade, dtype: int64
```

```
In [68]: sns.boxplot(x='grade',y='price',data=bm_final)
```

```
Out[68]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9b0fd5f8>
```




```
In [69]: bm_final['grade'] = bm_final['grade'].replace({1:0,3:0,4:1,5:1,6:1,7:2,8:3,9:3,10:4,11:4,12:4,13:4})
bm_final['grade'].unique()
```

```
Out[69]: array([2, 1, 3, 4, 'nan', 0], dtype=object)
```

```
In [70]: sns.boxplot(x='grade',y='price',data=bm_final)
```

```
Out[70]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9b607240>
```



```
In [71]: bm_price_grade = bm_final[['grade','price']]
bm_price_grade.groupby('grade').mean()
```

```
Out[71]:
```

	price
grade	
0	2.620000e+05
1	2.956212e+05
2	4.026813e+05
3	6.656492e+05
4	1.690680e+06
nan	5.331155e+05

```
In [72]: bm_final['grade'].value_counts()
```

```
Out[72]: 3      8834
         2      8098
         1      2062
         nan    1031
         4       444
         0         1
         Name: grade, dtype: int64
```

Since, grade 'nan' has a different mean price. And 'grade' is an ordinal variable. Assuming higher the grade, higher the price.

We can impute nan as a new category '2.5'

```
In [73]: bm_final['grade'] = bm_final['grade'].replace({'nan':2.5})
```

```
In [74]: bm_final['grade'].value_counts()
```

```
Out[74]: 3.0      8834
         2.0      8098
         1.0      2062
         2.5      1031
         4.0       444
         0.0         1
         Name: grade, dtype: int64
```

```
In [75]: bm_final.isnull().sum()
```

```
Out[75]: date          0
         bedrooms      0
         bathrooms     0
         sqft_living    0
         sqft_lot       0
         floors         0
         waterfront     0
         view           0
         condition      0
         grade          0
         sqft_above     1079
         sqft_basement  1030
         yr_built       971
         yr_renovated   1102
         zipcode        1060
         lat            1027
         long           1078
         sqft_living15  1041
         sqft_lot15     1059
         price          0
         sale_year      0
         sale_month     0
         dtype: int64
```

'sqft_above', 'sqft_basement'

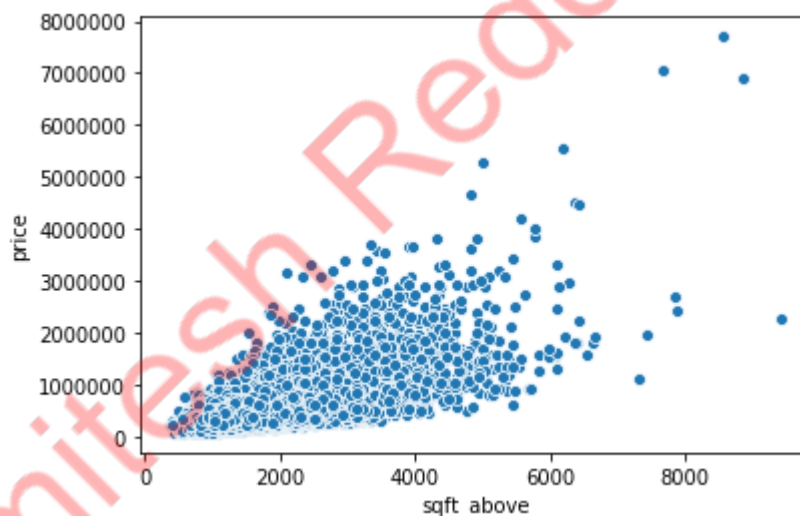
```
In [76]: bm_final['sqft_above'].describe()  
bm_final['sqft_basement'].describe()
```

```
Out[76]: count    19391.000000  
mean       1788.368934  
std        827.641928  
min         380.000000  
25%        1190.000000  
50%        1560.000000  
75%        2210.000000  
max         9410.000000  
Name: sqft_above, dtype: float64
```

```
Out[76]: count    19440.000000  
mean         290.629835  
std         442.046942  
min           0.000000  
25%           0.000000  
50%           0.000000  
75%          560.000000  
max         4820.000000  
Name: sqft_basement, dtype: float64
```


```
In [77]: sns.scatterplot(x='sqft_above',y = 'price',data=bm_final)
```

```
Out[77]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9b9d2be0>
```




```
In [78]: sns.scatterplot(x='sqft_basement', y='price', data=bm_final)
```

```
Out[78]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9ba48f98>
```



Filling the null values with mean and median

```
In [79]: bm_final['sqft_basement']=bm_final['sqft_basement'].fillna(bm_final['sqft_basement'].mean())  
bm_final['sqft_above']=bm_final['sqft_above'].fillna(bm_final['sqft_above'].mean())
```

```
In [80]: bm_final.isnull().sum()
```

```
Out[80]: date                0
bedrooms                   0
bathrooms                  0
sqft_living                0
sqft_lot                   0
floors                     0
waterfront                 0
view                       0
condition                  0
grade                      0
sqft_above                 0
sqft_basement              0
yr_built                   971
yr_renovated               1102
zipcode                   1060
lat                        1027
long                      1078
sqft_living15              1041
sqft_lot15                 1059
price                      0
sale_year                  0
sale_month                 0
dtype: int64
```

'yr_built'

Dropping the nan under yr_built column as the price range is same over the years and does not have an effect on the price except for some cases

```
In [81]: bm_final.yr_built.fillna('nan',inplace=True)
```


```
In [82]: bm_final['yr_built'].value_counts()
```

```
Out[82]: nan      971
        2014      503
        2005      410
        2006      398
        2004      385
        1977      382
        2003      366
        2007      363
        1968      350
        1978      333
        2008      327
        1967      314
        1979      308
        1959      302
        1990      287
        1962      283
        1954      282
        2001      279
        1987      267
        1989      256
        1969      255
        1955      244
        1988      244
        1947      243
        1999      241
        1994      239
        1976      233
        1950      233
        1963      231
        1966      224
        ...
        1945       82
        1906       82
        1909       81
        1919       80
        1908       78
        1900       78
        1923       71
        1916       71
        1912       70
        1905       69
        1921       68
        1911       68
        1937       61
        1907       59
        1931       55
        1915       54
        1917       51
        1913       49
        1914       48
        1938       46
        1904       44
        1903       43
        1936       39
        2015       36
```

```
1932      34
1933      29
1901      27
1902      25
1935      22
1934      21
Name: yr_built, Length: 117, dtype: int64
```

```
In [83]: sns.scatterplot(x='yr_built',y='price',data=bm_final)
```

```
Out[83]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9bab6a58>
```



Dropping the column 'yr_built' as the price for the houses built across the years have the same sales price

```
In [84]: bm_final = bm_final.drop(['yr_built'],axis=1)
```

```
In [85]: bm_final.shape  
bm_final.isnull().sum()
```

```
Out[85]: (20470, 21)
```

```
Out[85]: date                0  
bedrooms                    0  
bathrooms                   0  
sqft_living                  0  
sqft_lot                     0  
floors                       0  
waterfront                   0  
view                         0  
condition                    0  
grade                        0  
sqft_above                   0  
sqft_basement                0  
yr_renovated                1102  
zipcode                     1060  
lat                          1027  
long                         1078  
sqft_living15                1041  
sqft_lot15                   1059  
price                        0  
sale_year                    0  
sale_month                   0  
dtype: int64
```

'yr_renovated'

```
In [86]: bm_final.yr_renovated.fillna('nan',inplace=True)
bm_final['yr_renovated'].value_counts()
```

```
Out[86]: 0          18559
nan         1102
2014          81
2003          33
2013          31
2005          30
2000          29
2007          26
2004          24
1990          23
2006          22
1989          21
2009          21
2002          20
1998          19
2001          19
1987          18
1984          18
1991          18
1983          17
1994          17
2008          17
2010          17
1993          16
1997          15
1985          15
1992          14
1986          13
2015          13
1988          13
...
1968          5
1958          5
1981          5
1978          4
1965          4
1964          4
1973          4
1972          4
1969          3
1953          3
1955          3
1956          3
1974          3
1957          3
1960          3
1963          3
1940          2
1962          2
1967          2
1971          2
1976          2
1950          2
1946          2
```

```

1945      2
1934      1
1954      1
1944      1
1948      1
1951      1
1959      1

```

Name: yr_renovated, Length: 71, dtype: int64

```
In [87]: bm_final['yr_renovated'].unique()
```

```
Out[87]: array([0, 1991, 'nan', 2002, 2010, 1992, 2013, 2005, 2008, 2003, 1994,
        1984, 1954, 2014, 2011, 1974, 1999, 1983, 1990, 1988, 1957, 1977,
        1981, 1995, 1978, 2000, 1998, 1970, 1989, 2004, 1986, 2009, 2007,
        1987, 1973, 2006, 2001, 1980, 1971, 1945, 1979, 1997, 1950, 1948,
        2015, 2012, 1968, 1963, 1951, 1993, 1962, 1996, 1972, 1985, 1953,
        1955, 1982, 1956, 1969, 1940, 1946, 1975, 1958, 1964, 1976, 1959,
        1960, 1967, 1965, 1934, 1944], dtype=object)
```

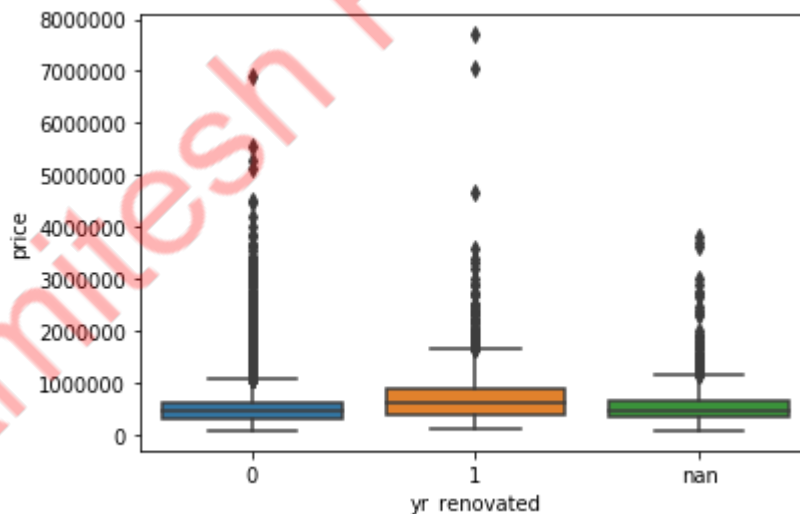
```
In [88]: bm_final.loc[((bm_final.yr_renovated != 0) & (bm_final.yr_renovated != 'nan')), 'yr_renovated']
```

```
In [89]: bm_final['yr_renovated'].value_counts()
```

```
Out[89]: 0      18559
        nan     1102
         1       809
        Name: yr_renovated, dtype: int64
```

```
In [90]: sns.boxplot(x='yr_renovated', y='price', data=bm_final)
```

```
Out[90]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9bae2860>
```



We can infer that the price of houses which are renovated is higher when compared to the non-renovated houses

Imputing nan with value '0' as they both have comparatively similar range

```
In [91]: bm_final['yr_renovated'] = bm_final['yr_renovated'].replace({'nan':0})
```

```
In [92]: bm_final['yr_renovated'].value_counts()  
bm_final.isnull().sum()
```

```
Out[92]: 0    19661  
        1     809  
        Name: yr_renovated, dtype: int64
```

```
Out[92]: date                0  
        bedrooms            0  
        bathrooms          0  
        sqft_living         0  
        sqft_lot            0  
        floors              0  
        waterfront          0  
        view                0  
        condition           0  
        grade               0  
        sqft_above          0  
        sqft_basement       0  
        yr_renovated        0  
        zipcode            1060  
        lat                 1027  
        long                1078  
        sqft_living15       1041  
        sqft_lot15          1059  
        price               0  
        sale_year           0  
        sale_month          0  
        dtype: int64
```

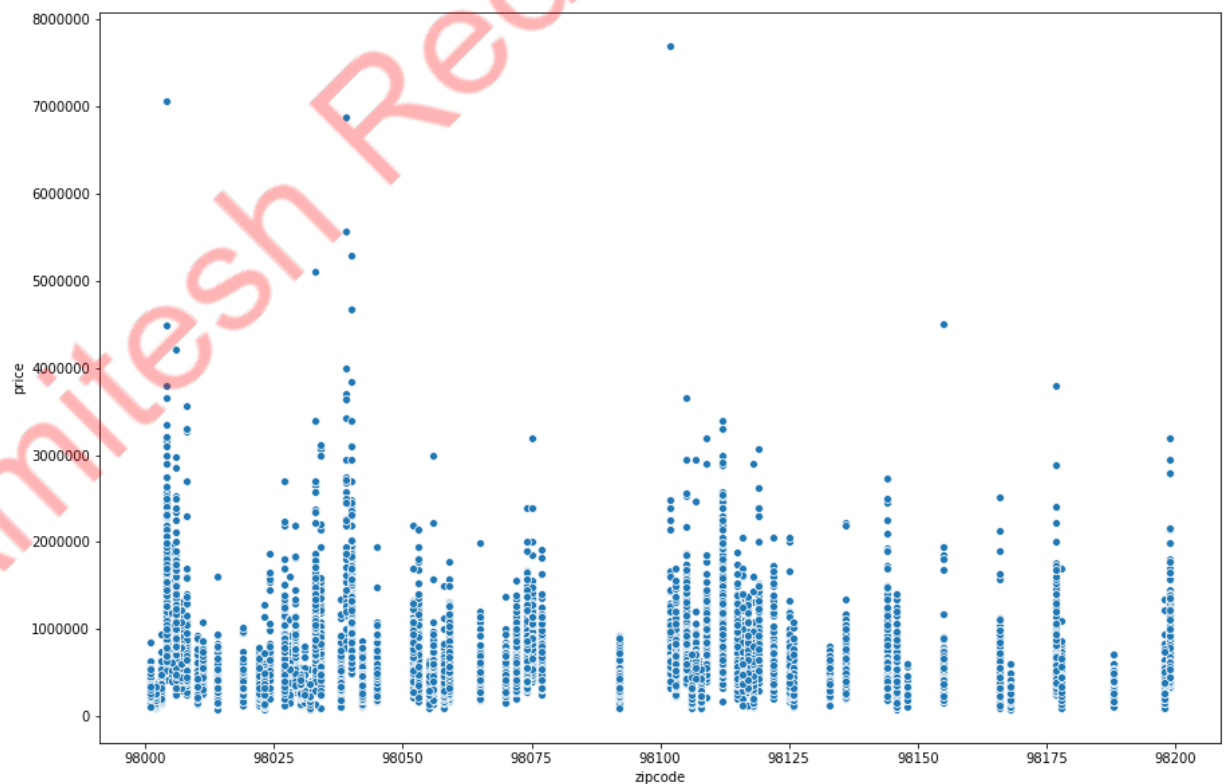
'zipcode'


```
In [93]: bm_final.zipcode.fillna('nan',inplace=True)
bm_final['zipcode'].value_counts()
```

```
Out[93]: nan      1060
98103      545
98038      537
98115      522
98052      516
98034      502
98117      500
98042      488
98118      457
98023      453
98006      446
98133      442
98059      419
98058      413
98155      402
98033      402
98074      397
98027      372
98125      361
98055      350
```

```
In [94]: a4_dims = (15,10)
fig, ax = plt.subplots(figsize=a4_dims)
sns.scatterplot(x='zipcode',y='price',data=bm_final)
```

```
Out[94]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9bb92eb8>
```



For the purpose of grouping the areas, we are calculating the price for each zipcode to form a group of zipcodes for easier analysis

```
In [95]: bm_final['zipcode'].value_counts()  
bm_price_floors = bm_final[['zipcode','price']]
```

```
Out[95]: nan      1060  
98103      545  
98038      537  
98115      522  
98052      516  
98034      502  
98117      500  
98042      488  
98118      457  
98023      453  
98006      446  
98133      442  
98059      419  
98058      413  
98155      402  
98033      402  
98074      397  
98027      372  
98125      361  
98056      360  
98053      358  
98126      323  
98001      321  
98075      318  
98144      310  
98106      306  
98092      305  
98116      298  
98029      297  
98004      294  
...  
98112      242  
98031      242  
98168      238  
98055      236  
98107      231  
98178      230  
98177      229  
98030      227  
98166      225  
98022      207  
98105      203  
98045      200  
98077      182  
98002      182  
98019      173  
98011      171  
98108      164  
98119      162  
98005      154  
98188      124  
98007      117  
98014      110  
98032      107
```

```
98070    102
98109    101
98102     91
98010     85
98024     69
98148     55
98039     44
```

Name: zipcode, Length: 71, dtype: int64

```
In [96]: data = bm_price_floors.groupby('zipcode').mean()
data = pd.DataFrame(data)
data['zipcode'] = data.index
data.columns
data.sort_values('price')
sns.scatterplot(x='zipcode',y='price',data=data)
```

Out[96]: Index(['price', 'zipcode'], dtype='object')

```
In [97]: data.loc[((data.price < 300000)), 'price'] = 1
data.loc[((data.price >= 300000) & (data.price <= 600000)), 'price'] = 2
data.loc[((data.price > 600000) & (data.price <= 900000)), 'price'] = 3
data.loc[(data.price >= 900000), 'price'] = 4
```

```
In [98]: data['price'].value_counts()
```

```
Out[98]: 2.0    36
3.0    20
1.0    10
4.0     5
Name: price, dtype: int64
```

```
In [99]: a = data.loc[data['price']==1]
a1 = a['zipcode'].unique()
a1
b = data.loc[data['price']==2]
b1 = b['zipcode'].unique()
b1
c = data.loc[data['price']==3]
c1 = c['zipcode'].unique()
c1
d = data.loc[data['price']==4]
d1 = d['zipcode'].unique()
d1
```

```
Out[99]: array([98001, 98002, 98003, 98023, 98030, 98031, 98032, 98148, 98168,
98188], dtype=object)
```

```
Out[99]: array([98010, 98011, 98014, 98019, 98022, 98024, 98028, 98034, 98038,
98042, 98045, 98055, 98056, 98058, 98059, 98065, 98070, 98072,
98092, 98103, 98106, 98107, 98108, 98117, 98118, 98125, 98126,
98133, 98136, 98144, 98146, 98155, 98166, 98178, 98198, 'nan'],
dtype=object)
```

```
Out[99]: array([98005, 98006, 98007, 98008, 98027, 98029, 98033, 98052, 98053,
98074, 98075, 98077, 98105, 98109, 98115, 98116, 98119, 98122,
98177, 98199], dtype=object)
```

```
Out[99]: array([98004, 98039, 98040, 98102, 98112], dtype=object)
```

We have identified the zipcodes to be combined

Dropping the null values and then proceeding with grouping of zipcodes

```
In [100]: bm_final = bm_final[bm_final.zipcode != 'nan']
bm_final['zipcode'].isnull().sum()
```

```
Out[100]: 0
```

```
In [101]: bm_final['zipcode'] = bm_final['zipcode'].astype(object)
```

```
In [102]: bm_final['zipcode'] = bm_final['zipcode'].replace(a1, '1')
bm_final['zipcode'] = bm_final['zipcode'].replace(b1, '2')
bm_final['zipcode'] = bm_final['zipcode'].replace(c1, '3')
bm_final['zipcode'] = bm_final['zipcode'].replace(d1, '4')
```

```
In [103]: bm_final['zipcode'].value_counts()
```

```
Out[103]: 2    10396
3     5880
1     2207
4      927
Name: zipcode, dtype: int64
```

```
In [104]: bm_final.shape  
          bm_final.isnull().sum()
```


```
Out[104]: (19410, 21)
```

```
Out[104]: date                0  
          bedrooms            0  
          bathrooms           0  
          sqft_living          0  
          sqft_lot             0  
          floors              0  
          waterfront           0  
          view                 0  
          condition            0  
          grade                0  
          sqft_above           0  
          sqft_basement        0  
          yr_renovated         0  
          zipcode              0  
          lat                  976  
          long                 1017  
          sqft_living15        984  
          sqft_lot15           1007  
          price                0  
          sale_year            0  
          sale_month           0  
          dtype: int64
```

'sqft_living15' and 'sqft_lot15'


```
In [105]: sns.scatterplot(x='sqft_living15',y='price',data=bm_final)
```

```
Out[105]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9bc3d6a0>
```



```
In [106]: sns.scatterplot(x='sqft_lot15',y='price',data=bm_final)
```

```
Out[106]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9bcea630>
```



```
In [107]: bm_final['sqft_living15']=bm_final['sqft_living15'].fillna(bm_final['sqft_living15'].mean)
bm_final['sqft_lot15']=bm_final['sqft_lot15'].fillna(bm_final['sqft_lot15'].mean)
```

```
In [108]: bm_final.isnull().sum()
```

```
Out[108]: date                0
bedrooms                    0
bathrooms                   0
sqft_living                  0
sqft_lot                     0
floors                       0
waterfront                   0
view                         0
condition                    0
grade                        0
sqft_above                   0
sqft_basement                 0
yr_renovated                 0
zipcode                      0
lat                          976
long                        1017
sqft_living15                0
sqft_lot15                   0
price                        0
sale_year                    0
sale_month                   0
dtype: int64
```

Dropping the date, lat, long, yr_built column

```
In [109]: bm_final = bm_final.drop(['date', 'lat', 'long'], axis=1)
```

Dataset after cleaning null values

```
In [110]: bm_final.shape
```

```
Out[110]: (19410, 18)
```

```
In [111]: bm_final.columns
```

```
Out[111]: Index(['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',
                'waterfront', 'view', 'condition', 'grade', 'sqft_above',
                'sqft_basement', 'yr_renovated', 'zipcode', 'sqft_living15',
                'sqft_lot15', 'price', 'sale_year', 'sale_month'],
                dtype='object')
```

```
In [112]: bm_final.head()
```

```
Out[112]:
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_abov
0	3	1.00	1180.0	5650.0	1.0	0	0	3	2.0	1180
1	3	2.25	2570.0	7242.0	2.0	0	0	3	2.0	2170
2	2	1.00	770.0	10000.0	1.0	0	0	3	1.0	770
3	4	3.00	1960.0	5000.0	1.0	0	0	5	2.0	1050
4	3	2.00	1680.0	8080.0	1.0	0	0	3	3.0	1680

```
In [113]: bm_final.dtypes
```


```
Out[113]: bedrooms      int64
bathrooms      float64
sqft_living     float64
sqft_lot        float64
floors          float64
waterfront      int64
view            int64
condition       int64
grade           float64
sqft_above      float64
sqft_basement   float64
yr_renovated    object
zipcode         object
sqft_living15   float64
sqft_lot15      float64
price           int64
sale_year       object
sale_month      object
dtype: object
```

Columns 'zipcode', 'sale_month' should have dtype as object (i.e.) categorical

```
In [114]: bm_final['yr_renovated'] = bm_final['yr_renovated'].astype(int)
bm_final['sale_year'] = bm_final['sale_year'].astype(int)
```

```
In [115]: bmcov = bm_final.corr()
a4_dims = (20,10)
fig, ax = plt.subplots(figsize=a4_dims)
sns.heatmap(bmcov,annot=True)
```

Out[115]: <matplotlib.axes._subplots.AxesSubplot at 0x1fd9bcea320>



There is no correlation of greater than 0.7 between price and any other variable

Performing One-hot vector encoding for the column 'zipcode'

```
In [116]: house_sales = pd.get_dummies(bm_final)
```



```
In [117]: house_sales.shape
house_sales.head(10)
```

```
Out[117]: (19410, 32)
```

```
Out[117]:
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above
0	3	1.00	1180.0	5650.0	1.0	0	0	3	2.0	118
1	3	2.25	2570.0	7242.0	2.0	0	0	3	2.0	217
2	2	1.00	770.0	10000.0	1.0	0	0	3	1.0	77
3	4	3.00	1960.0	5000.0	1.0	0	0	5	2.0	105
4	3	2.00	1680.0	8080.0	1.0	0	0	3	3.0	168
5	4	2.25	5420.0	7620.0	1.0	0	0	3	4.0	389
6	3	2.25	1715.0	6819.0	2.0	0	0	3	2.0	171
7	3	1.50	1060.0	9711.0	1.0	0	0	3	2.0	106
9	3	2.25	1890.0	6560.0	2.0	0	0	3	2.0	189
11	2	1.00	1160.0	6000.0	1.5	0	0	4	2.0	86

10 rows × 32 columns

Train and test split

```
In [118]: X = house_sales.drop(['price'],axis=1)
y = house_sales['price']

X_train_org, X_test_org, y_train, y_test = train_test_split(X, y, test_size=0.25)
y_train.shape
y_test.shape
X_train_org.shape
X_test_org.shape
```

```
Out[118]: (14557,)
```

```
Out[118]: (4853,)
```

```
Out[118]: (14557, 31)
```

```
Out[118]: (4853, 31)
```

Scaling

In our dataset, there are outliers present. Hence, we are using StandardScaler as the MinMax scaler is very sensitive to the presence of outliers.

```
In [119]: scaler = StandardScaler()
```

```
In [120]: X_train_scale = scaler.fit_transform(X_train_org)
```

```
In [121]: X_test_scale = scaler.transform(X_test_org)
```

Linear Regression

```
In [122]: from sklearn.model_selection import GridSearchCV
model = LinearRegression()
parameters = {'normalize':[True,False]}
grid_search_lr = GridSearchCV(model,parameters, cv=6, return_train_score=True)
grid_search_lr.fit(X_train_org, y_train)
print("Best parameters: {}".format(grid_search_lr.best_params_))
print("Best cross-validation score: {:.4f}".format(grid_search_lr.best_score_))
```

```
Out[122]: GridSearchCV(cv=6, error_score='raise-deprecating',
                        estimator=LinearRegression(copy_X=True, fit_intercept=True,
                                                    n_jobs=None, normalize=False),
                        iid='warn', n_jobs=None, param_grid={'normalize': [True, False]},
                        pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                        scoring=None, verbose=0)
```

Best parameters: {'normalize': False}

Best cross-validation score: 0.7038

```
In [123]: results = pd.DataFrame(grid_search_lr.cv_results_)
results
```

```
Out[123]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_normalize	params	spli
0	0.015868	0.003134	0.001048	0.001518	True	{'normalize': True}	
1	0.020165	0.001392	0.004478	0.000699	False	{'normalize': False}	

2 rows × 23 columns

```
In [124]: lreg = LinearRegression(normalize = True)
lreg.fit(X_train_org, y_train)
print(lreg.score(X_train_org, y_train))
print(lreg.score(X_test_org, y_test))
```

```
Out[124]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=True)

0.7070004856644225
0.7057053631724195
```

```
In [125]: from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score

kfold = KFold(n_splits=6)
print("Cross-validation scores:\n{}".format(cross_val_score(lreg , X_train_org, y_train, cv=kfold)
scores = cross_val_score(lreg , X_train_org, y_train, cv=kfold)
print(np.mean(scores))
```

Cross-validation scores:
[0.70963421 0.71608895 0.71988844 0.69750333 0.68327092 0.69657358]
0.7038265695949354

```
In [126]: #PLOT

%matplotlib inline
import matplotlib.pyplot as plt

X_train_array = X_train_org.to_numpy()

X_train_rm = X_train_array[:,2].reshape(-1,1)
lreg.fit(X_train_rm, y_train)
y_predict = lreg.predict(X_train_rm)


plt.plot(X_train_rm, y_predict, c = 'r')
plt.scatter(X_train_rm,y_train)
plt.xlabel('sqft_living')
```

Out[126]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=True)

Out[126]: [matplotlib.lines.Line2D at 0x1fd9e9eae0]

Out[126]: <matplotlib.collections.PathCollection at 0x1fd9e98f828>

Out[126]: Text(0.5, 0, 'sqft_living')



Linear Regression Result:

Best parameter: {'normalize': False}

Average Cross validation score: 0.7038

Test score: 0.7057

KNN Regression

```
In [127]: grid_parms_knn = {'n_neighbors':[1,5,10,15,20]}
```

```
In [128]: knn = KNeighborsRegressor()
grid_search_knn = GridSearchCV(knn, grid_parms_knn,cv=6,return_train_score=True,
grid_search_knn.fit(X_train_scale, y_train)
```

```
Out[128]: GridSearchCV(cv=6, error_score='raise-deprecating',
                        estimator=KNeighborsRegressor(algorithm='auto', leaf_size=30,
                                                    metric='minkowski',
                                                    metric_params=None, n_jobs=None,
                                                    n_neighbors=5, p=2,
                                                    weights='uniform'),
                        iid='warn', n_jobs=-1,
                        param_grid={'n_neighbors': [1, 5, 10, 15, 20]},
                        pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                        scoring=None, verbose=0)
```


```
In [129]: print("Best parameters: {}".format(grid_search_knn.best_params_))
print("Best cross-validation score: {:.4f}".format(grid_search_knn.best_score_))
pd.DataFrame(grid_search_knn.cv_results_)
```

Best parameters: {'n_neighbors': 10}
Best cross-validation score: 0.7051

```
Out[129]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_n_neighbors	params
0	0.355068	0.024509	4.756456	0.158985	1	{'n_neighbors': 1}
1	0.371450	0.033257	6.437256	0.259872	5	{'n_neighbors': 5}
2	0.394010	0.030371	7.297763	0.196907	10	{'n_neighbors': 10}
3	0.392970	0.030417	7.835041	0.268176	15	{'n_neighbors': 15}
4	0.392819	0.024051	8.136292	0.180613	20	{'n_neighbors': 20}

5 rows × 7 columns



```
In [130]: knn = KNeighborsRegressor(n_neighbors = 10)
knn.fit(X_train_scale, y_train)
print(knn.score(X_train_scale, y_train))
print(knn.score(X_test_scale, y_test))
```

```
Out[130]: KNeighborsRegressor(algorithm='auto', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=None, n_neighbors=10, p=2,
                             weights='uniform')
```

```
0.7596335157792412
0.7091584457948231
```

```
In [131]: from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score

kfold = KFold(n_splits=6)
print("Cross-validation scores:\n{}".format(cross_val_score(knn, X_train_scale,
scores = cross_val_score(knn, X_train_scale, y_train, cv=kfold)
print(np.mean(scores))
```

```
Cross-validation scores:
[0.7102932  0.71560329 0.71878694 0.71472988 0.6644096  0.70653242]
0.7050592204685852
```

```

In [132]: X_b = X_train_array[:50,2].reshape(-1,1)
           y_b = y_train[:50]

           knn_reg = KNeighborsRegressor(10)
           knn_reg.fit(X_b, y_b)

           X_new=np.linspace(X_b.min(), X_b.max(), 50).reshape(50, 1)
           y_predict = knn_reg.predict(X_new)

           plt.plot(X_new, y_predict, c = 'r')
           plt.scatter(X_b, y_b)
           plt.xlabel('sqft_living')

```

```

Out[132]: KNeighborsRegressor(algorithm='auto', leaf_size=30, metric='minkowski',
                               metric_params=None, n_jobs=None, n_neighbors=10, p=2,
                               weights='uniform')

```

```

Out[132]: [<matplotlib.lines.Line2D at 0x1fd9ea95a20>]

```

```

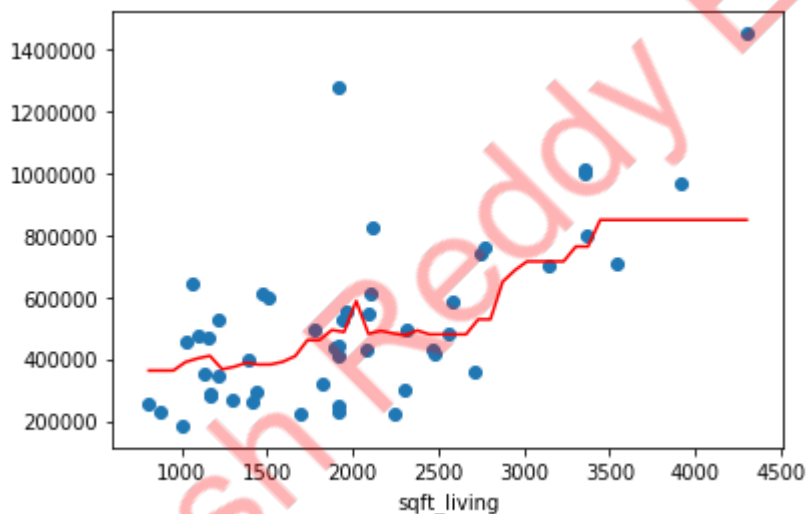
Out[132]: <matplotlib.collections.PathCollection at 0x1fd9ea95eb8>

```

```

Out[132]: Text(0.5, 0, 'sqft_living')

```



KNN Regression Result:

Best parameter: {n_neighbors: 10}

Average Cross validation score: 0.7050

Test score: 0.7091

Ridge Regression

```

In [133]: grid_parms_ridge = {'alpha': [0.01, 0.1, 1, 10, 100]}

```

```
In [134]: ridge = Ridge()
grid_search_ridge = GridSearchCV(estimator = ridge,param_grid = grid_params_ridge)
grid_search_ridge.fit(X_train_org, y_train)
print("Best parameters: {}".format(grid_search_ridge.best_params_))

print("Best cross-validation score: {:.4f}".format(grid_search_ridge.best_score_
```

```
Out[134]: GridSearchCV(cv=5, error_score='raise-deprecating',
                      estimator=Ridge(alpha=1.0, copy_X=True, fit_intercept=True,
                                     max_iter=None, normalize=False, random_state=None,
                                     solver='auto', tol=0.001),
                      iid='warn', n_jobs=-1,
                      param_grid={'alpha': [0.01, 0.1, 1, 10, 100]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                      scoring=None, verbose=0)
```

Best parameters: {'alpha': 1}
 Best cross-validation score: 0.7044

```
In [135]: ridge = Ridge(alpha = 1)
ridge.fit(X_train_org, y_train)
print(ridge.score(X_train_org, y_train))
print(ridge.score(X_test_org, y_test))
```

```
Out[135]: Ridge(alpha=1, copy_X=True, fit_intercept=True, max_iter=None, normalize=False,
                random_state=None, solver='auto', tol=0.001)

0.7069981831976242
0.7058582864583789
```

```
In [136]: from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score

kfold = KFold(n_splits=6)
print("Cross-validation scores:\n{}".format(cross_val_score(ridge , X_train_org,
scores = cross_val_score(ridge , X_train_org, y_train, cv=kfold)
print(np.mean(scores))
```


Cross-validation scores:
 [0.70978705 0.71599989 0.71998249 0.6977393 0.68300811 0.69661804]
 0.7038558124738913

```
In [137]: result_ridge = pd.DataFrame(grid_search_ridge.cv_results_)
result_ridge
```

```
Out[137]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_test_
0	0.016812	0.005477	0.003242	0.000977	0.01	{'alpha': 0.01}	0.7
1	0.018213	0.001392	0.003540	0.001892	0.1	{'alpha': 0.1}	0.7
2	0.018341	0.000349	0.002567	0.002119	1	{'alpha': 1}	0.7
3	0.016607	0.001885	0.004855	0.002448	10	{'alpha': 10}	0.7
4	0.017873	0.001684	0.001624	0.001988	100	{'alpha': 100}	0.7

5 rows × 21 columns




```
In [138]: import matplotlib.pyplot as plt
%matplotlib inline

plt.plot(range(result_ridge.shape[0]), result_ridge['mean_train_score'], label = 'mean train score')
plt.plot(range(result_ridge.shape[0]), result_ridge['mean_test_score'], label = 'mean test score')
plt.xticks(range(result_ridge.shape[0]), result_ridge['param_alpha'], rotation = 45)
plt.plot([grid_search_ridge.best_index_], result_ridge['mean_train_score'][grid_search_ridge.best_index_], label = 'best mean train score')
plt.plot([grid_search_ridge.best_index_], result_ridge['mean_test_score'][grid_search_ridge.best_index_], label = 'best mean test score')
plt.grid()
plt.legend()
plt.xlabel('Alpha')
```

Out[138]: [<matplotlib.lines.Line2D at 0x1fd9f88d9e8>]

Out[138]: [<matplotlib.lines.Line2D at 0x1fd9f8664a8>]


Out[138]: ([<matplotlib.axis.XTick at 0x1fd9ea587f0>, <matplotlib.axis.XTick at 0x1fd9ea56c88>, <matplotlib.axis.XTick at 0x1fd9ea565f8>, <matplotlib.axis.XTick at 0x1fd9f8983c8>, <matplotlib.axis.XTick at 0x1fd9f898898>], <a list of 5 Text xticklabel objects>)

Out[138]: [<matplotlib.lines.Line2D at 0x1fd9f8667f0>]

Out[138]: [<matplotlib.lines.Line2D at 0x1fd9f86d940>]

Out[138]: <matplotlib.legend.Legend at 0x1fd9f898240>

Out[138]: Text(0.5, 0, 'Alpha')



Ridge Regression Result:

Best parameter: {'alpha': 1}

Average Cross validation score: 0.7038

Test score: 0.7058

Lasso Regression

```
In [139]: grid_parms_lasso = {'alpha': [0.01, 0.1, 1, 10, 100]}
```

```
In [140]: lasso = Lasso()
grid_search_lasso = GridSearchCV(estimator = lasso, param_grid = grid_parms_lasso)
grid_search_lasso.fit(X_train_org, y_train)
print("Best parameters: {}".format(grid_search_lasso.best_params_))
print("Best cross-validation score: {:.4f}".format(grid_search_lasso.best_score_))
```

```
Out[140]: GridSearchCV(cv=5, error_score='raise-deprecating',
                      estimator=Lasso(alpha=1.0, copy_X=True, fit_intercept=True,
                                      max_iter=1000, normalize=False, positive=False,
                                      precompute=False, random_state=None,
                                      selection='cyclic', tol=0.0001, warm_start=False),
                      iid='warn', n_jobs=-1,
                      param_grid={'alpha': [0.01, 0.1, 1, 10, 100]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                      scoring=None, verbose=0)
```

```
Best parameters: {'alpha': 10}
Best cross-validation score: 0.7044
```

```
In [141]: lass = Lasso(alpha = 10)
lass.fit(X_train_org, y_train)
print(lass.score(X_train_org, y_train))
print(lass.score(X_test_org, y_test))
```

```
Out[141]: Lasso(alpha=10, copy_X=True, fit_intercept=True, max_iter=1000, normalize=False,
              positive=False, precompute=False, random_state=None, selection='cyclic',
              tol=0.0001, warm_start=False)
```

```
0.7069999591510197
0.7057416447269227
```

```
In [142]: from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score

kfold = KFold(n_splits=6)
print("Cross-validation scores:\n{}".format(cross_val_score(lass, X_train_org, y_train, cv=kfold)))
scores = cross_val_score(lass, X_train_org, y_train, cv=kfold)
print(np.mean(scores))
```


```
Cross-validation scores:
[0.7096228  0.7160749  0.71994196  0.69756932  0.68323548  0.69658378]
0.7038380413758031
```

```
In [143]: result_lasso = pd.DataFrame(grid_search_lasso.cv_results_)
result_lasso
```

```
Out[143]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_test_
0	3.353209	0.186871	0.007256	0.001912	0.01	{'alpha': 0.01}	0.7
1	3.576854	0.068750	0.007105	0.003076	0.1	{'alpha': 0.1}	0.7
2	1.758889	0.584525	0.005954	0.002353	1	{'alpha': 1}	0.7
3	0.592850	0.103912	0.009030	0.011775	10	{'alpha': 10}	0.7
4	0.282584	0.028904	0.006506	0.002697	100	{'alpha': 100}	0.7

5 rows × 21 columns



In [144]: %matplotlib inline

```
plt.plot(range(result_lasso.shape[0]), result_lasso['mean_train_score'], label =
plt.plot(range(result_lasso.shape[0]), result_lasso['mean_test_score'], label =
plt.xticks(range(result_lasso.shape[0]), result_lasso['param_alpha'], rotation =
plt.plot([grid_search_lasso.best_index_], result_lasso['mean_train_score'][grid_s
plt.plot([grid_search_lasso.best_index_], result_lasso['mean_test_score'][grid_s
plt.grid()
plt.legend()
plt.xlabel('Alpha')
```

Out[144]: [<matplotlib.lines.Line2D at 0x1fd9f90f518>]

Out[144]: [<matplotlib.lines.Line2D at 0x1fd9f8ddc50>]

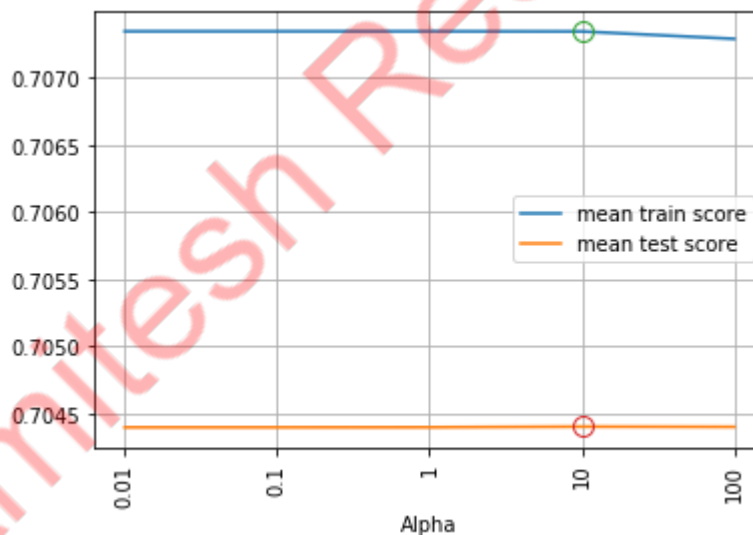
Out[144]: ([<matplotlib.axis.XTick at 0x1fd9f8c6c88>,
<matplotlib.axis.XTick at 0x1fd9f8ac8d0>,
<matplotlib.axis.XTick at 0x1fd9f8d4550>,
<matplotlib.axis.XTick at 0x1fd9f90feb8>,
<matplotlib.axis.XTick at 0x1fd9f9173c8>],
<a list of 5 Text xticklabel objects>)

Out[144]: [<matplotlib.lines.Line2D at 0x1fd9f90fb70>]

Out[144]: [<matplotlib.lines.Line2D at 0x1fd9f8d4048>]

Out[144]: <matplotlib.legend.Legend at 0x1fd9f9200f0>

Out[144]: Text(0.5, 0, 'Alpha')



Lasso Regression Result:

Best parameter: {'alpha': 100}

Average Cross validation score: 0.7038

Test score: 0.7057

Polynomial Regression

```
In [145]: from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.pipeline import make_pipeline

def PolynomialRegression(degree=2, **kwargs):
    return make_pipeline(PolynomialFeatures(degree),
                          LinearRegression(**kwargs))
```

```
In [146]: param_grid_poly = {'polynomialfeatures__degree': np.arange(3)}

grid_poly = GridSearchCV(PolynomialRegression(), param_grid_poly, return_train_score=True)
```

```
In [147]: grid_poly.fit(X_train_org, y_train)
```

```
Out[147]: GridSearchCV(cv=5, error_score='raise-deprecating',
                        estimator=Pipeline(memory=None,
                                           steps=[('polynomialfeatures',
                                                    PolynomialFeatures(degree=2,
                                                                           include_bias=True,
                                                                           interaction_only=False,
                                                                           order='C')),
                                                  ('linearregression',
                                                   LinearRegression(copy_X=True,
                                                                      fit_intercept=True,
                                                                      n_jobs=None,
                                                                      normalize=False))],
                                           verbose=False),
                        iid='warn', n_jobs=-1,
                        param_grid={'polynomialfeatures__degree': array([0, 1, 2])},
                        pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                        scoring=None, verbose=0)
```

```
In [148]: print("Best parameters: {}".format(grid_poly.best_params_))
print("Best cross-validation score: {:.4f}".format(grid_poly.best_score_))
```

```
Best parameters: {'polynomialfeatures__degree': 2}
Best cross-validation score: 0.7915
```

```
In [149]: pol = PolynomialFeatures(degree = 2)
X_pol = pol.fit_transform(X_train_org)
Xt_pol = pol.fit_transform(X_test_org)
pol_reg = LinearRegression()
pol_reg.fit(X_pol,y_train)
print(pol_reg.score(X_pol, y_train))
print(pol_reg.score(Xt_pol, y_test))
```

```
Out[149]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

0.8278011097032606
0.7942424473052565
```

```
In [150]: from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score

kfold = KFold(n_splits=6)
print("Cross-validation scores:\n{}".format(cross_val_score(pol_reg , X_pol, y_train, cv=kfold)
scores = cross_val_score(pol_reg , X_pol, y_train, cv=kfold)
print(np.mean(scores))
```


```
Cross-validation scores:
[0.78502321 0.7979335  0.78108883 0.79162343 0.809837  0.76740372]
0.7888182811260895
```

```
In [151]: result_poly = pd.DataFrame(grid_poly.cv_results_)
result_poly
```

```
Out[151]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_polynomialfeatures__degree
0	0.012437	0.002573	0.002110	0.001110	0
1	0.048420	0.004095	0.007297	0.002061	1
2	1.099465	0.046975	0.064668	0.014644	2

```
3 rows × 6 columns
```



In [152]:

```
plt.plot(range(result_poly.shape[0]), result_poly['mean_train_score'], label = 'mean train score')
plt.plot(range(result_poly.shape[0]), result_poly['mean_test_score'], label = 'mean test score')
plt.xticks(range(result_poly.shape[0]), result_poly['param_polynomialfeatures_degree'])
plt.plot([grid_poly.best_index_], result_poly['mean_train_score'][grid_poly.best_index_])
plt.plot([grid_poly.best_index_], result_poly['mean_test_score'][grid_poly.best_index_])
plt.grid()
plt.xlabel('Degree')
plt.legend()
```

Out[152]: [<matplotlib.lines.Line2D at 0x1fd9f990dd8>]

Out[152]: [<matplotlib.lines.Line2D at 0x1fd9f99c240>]

Out[152]: ([<matplotlib.axis.XTick at 0x1fd9f9783c8>,
 <matplotlib.axis.XTick at 0x1fd9f976cc0>,
 <matplotlib.axis.XTick at 0x1fd9f8ddbe0>],
 <a list of 3 Text xticklabel objects>)

Out[152]: [<matplotlib.lines.Line2D at 0x1fd9f99c630>]

Out[152]: [<matplotlib.lines.Line2D at 0x1fd9f8dd7b8>]

Out[152]: Text(0.5, 0, 'Degree')

Out[152]: <matplotlib.legend.Legend at 0x1fd9f99cdd8>



Polynomial Regression Result:

Best parameters: {'polynomialfeatures__degree': 2}

Average Cross validation score: 0.7888

Test score: 0.7942

Linear (Simple) SVR

```
In [153]: grid_params_svr1 = {'C': [0.01, 0.1, 1, 10, 100], 'epsilon' : [0.01, 0.1, 1, 10, 100]}
```

```
In [154]: linearsvr = LinearSVR()
grid_svr1 = GridSearchCV(estimator = linearsvr, param_grid = grid_params_svr1, return_train_score=True)
```

```
In [155]: grid_svr1.fit(X_train_scale, y_train)
```

```
Out[155]: GridSearchCV(cv=10, error_score='raise-deprecating',
                        estimator=LinearSVR(C=1.0, dual=True, epsilon=0.0,
                                             fit_intercept=True, intercept_scaling=1.0,
                                             loss='epsilon_insensitive', max_iter=1000,
                                             random_state=None, tol=0.0001, verbose=0),
                        iid='warn', n_jobs=-1,
                        param_grid={'C': [0.01, 0.1, 1, 10, 100],
                                    'epsilon': [0.01, 0.1, 1, 10, 100]},
                        pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                        scoring=None, verbose=0)
```

```
In [156]: print("Best parameters: {}".format(grid_svr1.best_params_))
print("Best cross-validation score: {:.4f}".format(grid_svr1.best_score_))
```

Best parameters: {'C': 100, 'epsilon': 10}

Best cross-validation score: 0.5549

```
In [157]: lsvr = LinearSVR(C = 100, epsilon = 1)

lsvr.fit(X_train_scale, y_train)

print(lsvr.score(X_train_scale, y_train))
print(lsvr.score(X_test_scale, y_test))
```

```
Out[157]: LinearSVR(C=100, dual=True, epsilon=1, fit_intercept=True,
                    intercept_scaling=1.0, loss='epsilon_insensitive', max_iter=1000,
                    random_state=None, tol=0.0001, verbose=0)
```

0.5676259465714262

0.5714589745919514


```
In [158]: from sklearn.model_selection import KFold
          from sklearn.model_selection import cross_val_score

          kfold = KFold(n_splits=10)
          print("Cross-validation scores:\n{}".format(cross_val_score(lsvr , X_train_scale
          scores = cross_val_score(lsvr, X_train_scale, y_train, cv=kfold)
          print(np.mean(scores))
```

Cross-validation scores:

```
[0.5438185  0.57203006 0.56323625 0.555753   0.58539792 0.58296943
 0.55636097 0.51814406 0.52171033 0.54886689]
0.5548712794044903
```

Amitesh Reddy Enugala


```
In [159]: result_linearsvr = pd.DataFrame(grid_svr1.cv_results_)
result_linearsvr
```

```
Out[159]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_epsilon	param
0	0.064522	0.005136	0.002878	0.001977	0.01	0.01	{'C': 0.01, 'epsilon': 0.01}
1	0.065998	0.004664	0.002690	0.002005	0.01	0.1	{'C': 0.01, 'epsilon': 0.1}
2	0.060206	0.002599	0.002000	0.001946	0.01	1	{'C': 0.01, 'epsilon': 1}
3	0.064692	0.005485	0.001503	0.001773	0.01	10	{'C': 0.01, 'epsilon': 10}
4	0.063528	0.003081	0.001998	0.003066	0.01	100	{'C': 0.01, 'epsilon': 100}
5	0.061185	0.006544	0.002440	0.001994	0.1	0.01	{'C': 0.1, 'epsilon': 0.01}
6	0.065425	0.009028	0.001870	0.001979	0.1	0.1	{'C': 0.1, 'epsilon': 0.1}
7	0.066129	0.007779	0.001425	0.002004	0.1	1	{'C': 0.1, 'epsilon': 1}
8	0.064609	0.012102	0.001606	0.002114	0.1	10	{'C': 0.1, 'epsilon': 10}
9	0.058471	0.006269	0.001238	0.001697	0.1	100	{'C': 0.1, 'epsilon': 100}
10	0.051752	0.004463	0.002725	0.001898	1	0.01	{'C': 1, 'epsilon': 0.01}
11	0.054823	0.006824	0.001762	0.001708	1	0.1	{'C': 1, 'epsilon': 0.1}
12	0.068827	0.004589	0.002457	0.002999	1	1	{'C': 1, 'epsilon': 1}
13	0.066519	0.004266	0.002116	0.002260	1	10	{'C': 1, 'epsilon': 10}
14	0.068104	0.005309	0.002134	0.002108	1	100	{'C': 1, 'epsilon': 100}

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_epsilon	param
15	0.075731	0.006666	0.002213	0.001950	10	0.01	{'C': 10, 'epsilon': 0.01}
16	0.079332	0.012926	0.002581	0.001618	10	0.1	{'C': 10, 'epsilon': 0.1}
17	0.076175	0.011934	0.001921	0.002220	10	1	{'C': 10, 'epsilon': 1}
18	0.079074	0.013626	0.002255	0.001957	10	10	{'C': 10, 'epsilon': 10}
19	0.069037	0.009133	0.002244	0.002093	10	100	{'C': 10, 'epsilon': 100}
20	0.101198	0.016539	0.002969	0.002135	100	0.01	{'C': 100, 'epsilon': 0.01}
21	0.094971	0.011703	0.001738	0.001974	100	0.1	{'C': 100, 'epsilon': 0.1}
22	0.089644	0.012913	0.001816	0.002303	100	1	{'C': 100, 'epsilon': 1}
23	0.086846	0.007896	0.002433	0.002096	100	10	{'C': 100, 'epsilon': 10}
24	0.086774	0.016526	0.001602	0.001690	100	100	{'C': 100, 'epsilon': 100}

25 rows × 32 columns



```
In [160]: plt.plot(range(result_linearsvr.shape[0]), result_linearsvr['mean_train_score'],
plt.plot(range(result_linearsvr.shape[0]), result_linearsvr['mean_test_score'],
plt.xticks(range(result_linearsvr.shape[0]), result_linearsvr['param_C'], rotation=45)
plt.plot([grid_svr1.best_index_], result_linearsvr['mean_train_score'][grid_svr1.best_index_],
plt.plot([grid_svr1.best_index_], result_linearsvr['mean_test_score'][grid_svr1.best_index_],
plt.grid()
plt.legend()
plt.xlabel('Alpha')
```

Out[160]: [<matplotlib.lines.Line2D at 0x1fd9fa47550>]

Out[160]: [<matplotlib.lines.Line2D at 0x1fd9fa47908>]

Out[160]: ([<matplotlib.axis.XTick at 0x1fd9f9fb9b0>, <matplotlib.axis.XTick at 0x1fd9f9fbcc0>, <matplotlib.axis.XTick at 0x1fd9f9326d8>, <matplotlib.axis.XTick at 0x1fd9fa47f60>, <matplotlib.axis.XTick at 0x1fd9fa51400>, <matplotlib.axis.XTick at 0x1fd9fa51940>, <matplotlib.axis.XTick at 0x1fd9fa51e10>, <matplotlib.axis.XTick at 0x1fd9fa57320>, <matplotlib.axis.XTick at 0x1fd9fa577f0>, <matplotlib.axis.XTick at 0x1fd9fa57cc0>, <matplotlib.axis.XTick at 0x1fd9fa5e1d0>, <matplotlib.axis.XTick at 0x1fd9fa5e710>, <matplotlib.axis.XTick at 0x1fd9fa57470>, <matplotlib.axis.XTick at 0x1fd9fa51a20>, <matplotlib.axis.XTick at 0x1fd9fa5e8d0>, <matplotlib.axis.XTick at 0x1fd9fa65198>, <matplotlib.axis.XTick at 0x1fd9fa656d8>, <matplotlib.axis.XTick at 0x1fd9fa65c50>, <matplotlib.axis.XTick at 0x1fd9fa6a208>, <matplotlib.axis.XTick at 0x1fd9fa6a780>, <matplotlib.axis.XTick at 0x1fd9fa6acf8>, <matplotlib.axis.XTick at 0x1fd9fa722b0>, <matplotlib.axis.XTick at 0x1fd9fa6add8>, <matplotlib.axis.XTick at 0x1fd9fa650b8>, <matplotlib.axis.XTick at 0x1fd9fa72208>], <a list of 25 Text xticklabel objects>)

Out[160]: [<matplotlib.lines.Line2D at 0x1fd9fa47a20>]

Out[160]: [<matplotlib.lines.Line2D at 0x1fd9fa78dd8>]

Out[160]: <matplotlib.legend.Legend at 0x1fd9fa47898>

Out[160]: Text(0.5, 0, 'Alpha')


```
In [164]: print("Best parameters: {}".format(grid_svr_linear.best_params_))
print("Best cross-validation score: {:.4f}".format(grid_svr_linear.best_score_))
```

```
Best parameters: {'C': 100}
Best cross-validation score: 0.6314
```

```
In [165]: svr = SVR(kernel = 'linear', C = 100)
```

```
    #train the model
svr.fit(X_train_scale, y_train)

    #evaluate the model
print(svr.score(X_train_scale, y_train))
print(svr.score(X_test_scale, y_test))
```

```
Out[165]: SVR(C=100, cache_size=200, coef0=0.0, degree=3, epsilon=0.1,
gamma='auto_deprecated', kernel='linear', max_iter=-1, shrinking=True,
tol=0.001, verbose=False)
```

```
0.6348349244684914
0.6353862043822313
```

```
In [166]: from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score

kfold = KFold(n_splits=6)
print("Cross-validation scores:\n{}".format(cross_val_score(svr , X_train_scale,
scores = cross_val_score(svr , X_train_scale, y_train, cv=kfold)
print(np.mean(scores))
```

```
Cross-validation scores:
[0.63422537 0.63669211 0.64669841 0.64797949 0.5816611  0.64105024]
0.6313844530856692
```

```
In [167]: result_svr_linear = pd.DataFrame(grid_svr_linear.cv_results_)
result_svr_linear
```

```
Out[167]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params	split0_test_sco
0	27.320662	0.404486	3.332883	0.072339	0.01	{'C': 0.01}	-0.05474
1	27.429692	0.486008	3.346784	0.068195	0.1	{'C': 0.1}	-0.04141
2	26.220203	0.255432	3.288037	0.050573	1	{'C': 1}	0.07010
3	27.259319	1.066409	3.357370	0.044849	10	{'C': 10}	0.45786
4	15.662856	0.237806	1.837903	0.035668	100	{'C': 100}	0.63422

```
5 rows × 23 columns
```

```
In [168]: plt.plot(range(result_svr_linear.shape[0]), result_svr_linear['mean_train_score'])
plt.plot(range(result_svr_linear.shape[0]), result_svr_linear['mean_test_score'])
plt.xticks(range(result_svr_linear.shape[0]), result_svr_linear['param_C'], rotation=45)
plt.plot([grid_svr_linear.best_index_], result_svr_linear['mean_train_score'][grid_svr_linear.best_index_])
plt.plot([grid_svr_linear.best_index_], result_svr_linear['mean_test_score'][grid_svr_linear.best_index_])
plt.grid()
plt.legend()
plt.xlabel('C')
```

Out[168]: [<matplotlib.lines.Line2D at 0x1fd9ff6fe48>]

Out[168]: [<matplotlib.lines.Line2D at 0x1fd9ff77240>]

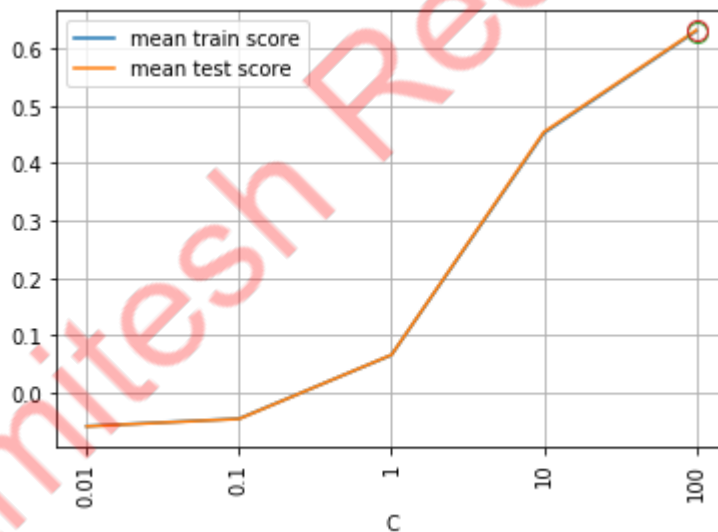
Out[168]: ([<matplotlib.axis.XTick at 0x1fd9ff56470>, <matplotlib.axis.XTick at 0x1fd9ff9e8a58>, <matplotlib.axis.XTick at 0x1fd9ff9f0da0>, <matplotlib.axis.XTick at 0x1fd9ff77898>, <matplotlib.axis.XTick at 0x1fd9ff77d68>], <a list of 5 Text xticklabel objects>)

Out[168]: [<matplotlib.lines.Line2D at 0x1fd9ff77550>]

Out[168]: [<matplotlib.lines.Line2D at 0x1fd9ff9f0940>]

Out[168]: <matplotlib.legend.Legend at 0x1fd9ff7ba90>

Out[168]: Text(0.5, 0, 'C')



SVR with Kernel as 'Linear' Result:

Best parameters: {'C': 100}

Average Cross validation score: 0.6313

Test score: 0.6353

SVR with kernel 'Poly'

```
In [169]: grid_params_svrp = {'C': [1, 10, 100], 'degree': [1, 3]}
```

```
In [170]: svr_poly = SVR(kernel='poly')
grid_svr_poly = GridSearchCV(estimator = svr_poly, param_grid = grid_params_svrp, n
```

```
In [171]: grid_svr_poly.fit(X_train_scale, y_train)
```

```
Out[171]: GridSearchCV(cv=3, error_score='raise-deprecating',
                        estimator=SVR(C=1.0, cache_size=200, coef0=0.0, degree=3,
                                      epsilon=0.1, gamma='auto_deprecated', kernel='poly',
                                      max_iter=-1, shrinking=True, tol=0.001,
                                      verbose=False),
                        iid='warn', n_jobs=-1,
                        param_grid={'C': [1, 10, 100], 'degree': [1, 3]},
                        pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                        scoring=None, verbose=0)
```

```
In [172]: print("Best parameters: {}".format(grid_svr_poly.best_params_))
print("Best cross-validation score: {:.4f}".format(grid_svr_poly.best_score_))
pd.DataFrame(grid_svr_poly.cv_results_)
```

Best parameters: {'C': 100, 'degree': 1}

Best cross-validation score: 0.2062

```
Out[172]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_degree	params
0	14.498907	0.041193	4.631024	0.126628	1	1	{'C': 1, 'degree': 1}
1	14.305654	0.086732	4.967052	0.100506	1	3	{'C': 1, 'degree': 3}
2	15.982145	0.304590	4.600349	0.302120	10	1	{'C': 10, 'degree': 1}
3	15.743284	0.281034	4.793856	0.259102	10	3	{'C': 10, 'degree': 3}
4	12.637298	0.144064	3.299894	0.060482	100	1	{'C': 100, 'degree': 1}
5	12.474434	0.247785	3.334822	0.039191	100	3	{'C': 100, 'degree': 3}


```
In [174]: svr_p = SVR(kernel='poly',C=100,degree = 1)
svr_p.fit(X_train_scale, y_train)
svr_p.score(X_train_scale, y_train)
svr_p.score(X_test_scale, y_test)
```

```
Out[174]: SVR(C=100, cache_size=200, coef0=0.0, degree=1, epsilon=0.1,
gamma='auto_deprecated', kernel='poly', max_iter=-1, shrinking=True,
tol=0.001, verbose=False)
```

```
Out[174]: 0.2837454148754319
```

```
Out[174]: 0.29004952647656823
```

Amitesh Reddy Enugala

```
In [175]: from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
#scores = cross_val_score(logreg, iris.data, iris.target)
kfold = KFold(n_splits=6)
print("Cross-validation scores:\n{}".format(cross_val_score(svr_p, X_train_scale,
scores = cross_val_score(svr_p, X_train_scale, y_train, cv=kfold)
print(np.mean(scores))
```

C:\Users\sures\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

C:\Users\sures\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

C:\Users\sures\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

C:\Users\sures\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

C:\Users\sures\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

C:\Users\sures\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

Cross-validation scores:

[0.25259198 0.25069119 0.25842774 0.26232613 0.21306053 0.25761964]

C:\Users\sures\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

C:\Users\sures\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

C:\Users\sures\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

C:\Users\sures\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

C:\Users\sures\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

C:\Users\sures\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

0.24911953373197618

In [176]: result_svr_poly = pd.DataFrame(grid_svr_poly.cv_results_)
result_svr_poly

Out[176]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_degree	param_
0	14.498907	0.041193	4.631024	0.126628	1	1	{'C': 1, 'degree': 1}
1	14.305654	0.086732	4.967052	0.100506	1	3	{'C': 1, 'degree': 3}
2	15.982145	0.304590	4.600349	0.302120	10	1	{'C': 10, 'degree': 1}
3	15.743284	0.281034	4.793856	0.259102	10	3	{'C': 10, 'degree': 3}
4	12.637298	0.144064	3.299894	0.060482	100	1	{'C': 100, 'degree': 1}

```
In [177]: plt.plot(range(result_svr_poly.shape[0]), result_svr_poly['mean_train_score'], label='mean train score')
plt.plot(range(result_svr_poly.shape[0]), result_svr_poly['mean_test_score'], label='mean test score')
plt.xticks(range(result_svr_poly.shape[0]), result_svr_poly['param_C'], rotation=45)
plt.plot([grid_svr_poly.best_index_], result_svr_poly['mean_train_score'][grid_svr_poly.best_index_], label='best train score')
plt.plot([grid_svr_poly.best_index_], result_svr_poly['mean_test_score'][grid_svr_poly.best_index_], label='best test score')
plt.grid()
plt.legend()
```

Out[177]: [<matplotlib.lines.Line2D at 0x1fd9fff49b0>]

Out[177]: [<matplotlib.lines.Line2D at 0x1fd9fff4d68>]

Out[177]: ([<matplotlib.axis.XTick at 0x1fd9ffd6048>, <matplotlib.axis.XTick at 0x1fd9ffd1908>, <matplotlib.axis.XTick at 0x1fd9ffd14e0>, <matplotlib.axis.XTick at 0x1fda0000400>, <matplotlib.axis.XTick at 0x1fda00008d0>, <matplotlib.axis.XTick at 0x1fda0000da0>], <a list of 6 Text xticklabel objects>)

Out[177]: [<matplotlib.lines.Line2D at 0x1fd9fff4e80>]

Out[177]: [<matplotlib.lines.Line2D at 0x1fd9fa5ebe0>]

Out[177]: <matplotlib.legend.Legend at 0x1fda000aba8>



SVR with Kernel as 'Poly' Result:

Best parameters: {'C': 100, 'degree': 1}

Average Cross validation score: 0.2491

Test score: 0.2900

SVR with kernel 'rbf'

```
In [178]: grid_parms_rbf = {'C': [0.1, 1, 10, 100], 'gamma': [0.1, 1, 10, 100]}
```

```
In [179]: svr_rbf = SVR(kernel='rbf')
grid_svr_rbf = GridSearchCV(estimator = svr_rbf, param_grid = grid_parms_rbf, return_train_score=True)
```

```
In [180]: grid_svr_rbf.fit(X_train_scale, y_train)
```

```
Out[180]: GridSearchCV(cv=3, error_score='raise-deprecating',
                      estimator=SVR(C=1.0, cache_size=200, coef0=0.0, degree=3,
                                   epsilon=0.1, gamma='auto_deprecated', kernel='rbf',
                                   max_iter=-1, shrinking=True, tol=0.001,
                                   verbose=False),
                      iid='warn', n_jobs=-1,
                      param_grid={'C': [0.1, 1, 10, 100], 'gamma': [0.1, 1, 10, 100]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                      scoring=None, verbose=0)
```

```
In [181]: print("Best parameters: {}".format(grid_svr_rbf.best_params_))
print("Best cross-validation score: {:.4f}".format(grid_svr_rbf.best_score_))
```

```
Best parameters: {'C': 100, 'gamma': 0.1}
Best cross-validation score: -0.0412
```

```
In [193]: svr_rbf = SVR(kernel='rbf', C=100, gamma=0.1)
svr_rbf.fit(X_train_scale, y_train)
svr_rbf.score(X_train_scale, y_train)
svr_rbf.score(X_test_scale, y_test)
```

```
Out[193]: SVR(C=100, cache_size=200, coef0=0.0, degree=3, epsilon=0.1, gamma=0.1,
             kernel='rbf', max_iter=-1, shrinking=True, tol=0.001, verbose=False)
```

```
Out[193]: -0.032844858673138466
```

```
Out[193]: -0.03450470646512982
```

```
In [183]: from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
kfold = KFold(n_splits=6)
print("Cross-validation scores:\n{}".format(cross_val_score(svr_rbf, X_train_scale, y_train, cv=kfold)))
scores = cross_val_score(svr_rbf, X_train_scale, y_train, cv=kfold)
print(np.mean(scores))
```

```
Cross-validation scores:
[-0.03381093 -0.04146735 -0.03190031 -0.03328871 -0.03862546 -0.04433172]
-0.03723741473126533
```

```
In [184]: result_rbf = pd.DataFrame(grid_svr_rbf.cv_results_)
result_rbf
```

```
Out[184]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_gamma	param
0	11.330334	0.125268	3.097708	0.139693	0.1	0.1	{'C': 0.1, 'gamma': 0.1}
1	12.243555	0.381897	3.180818	0.358904	0.1	1	{'C': 0.1, 'gamma': 1}
2	12.032121	0.478521	4.028549	0.353016	0.1	10	{'C': 0.1, 'gamma': 10}
3	13.576654	0.615906	5.750605	0.626988	0.1	100	{'C': 0.1, 'gamma': 100}
4	10.657468	0.084945	3.655548	0.062237	1	0.1	{'C': 1, 'gamma': 0.1}
5	11.766167	0.103654	3.543514	0.137770	1	1	{'C': 1, 'gamma': 1}
6	11.654134	0.049942	4.158866	0.026475	1	10	{'C': 1, 'gamma': 10}
7	13.023799	0.016312	5.219028	0.035598	1	100	{'C': 1, 'gamma': 100}
8	10.410131	0.081175	3.290524	0.038172	10	0.1	{'C': 10, 'gamma': 0.1}
9	11.356603	0.081094	3.384267	0.019435	10	1	{'C': 10, 'gamma': 1}
10	11.243234	0.070768	4.266910	0.031012	10	10	{'C': 10, 'gamma': 10}
11	13.335964	0.142453	5.280530	0.037150	10	100	{'C': 10, 'gamma': 100}
12	10.820697	0.037753	3.319779	0.130962	100	0.1	{'C': 100, 'gamma': 0.1}
13	11.597286	0.150764	3.535534	0.116806	100	1	{'C': 100, 'gamma': 1}
14	11.307395	0.328424	3.757275	0.191096	100	10	{'C': 100, 'gamma': 10}
15	9.973311	0.485107	3.137954	0.016946	100	100	{'C': 100, 'gamma': 100}

```
In [185]: plt.plot(range(result_rbf.shape[0]), result_rbf['mean_train_score'], label = 'me
plt.plot(range(result_rbf.shape[0]), result_rbf['mean_test_score'], label = 'mean
plt.xticks(range(result_svr_poly.shape[0]), result_rbf['param_C'], rotation = 90
plt.plot([grid_svr_rbf.best_index_], result_rbf['mean_train_score'][grid_svr_rbf
plt.plot([grid_svr_rbf.best_index_], result_rbf['mean_test_score'][grid_svr_rbf.
plt.grid()
plt.legend()
```

```
Out[185]: []
```


```
Out[185]: []
```

```
Out[185]: ([<matplotlib.axis.XTick at 0x1fda0046518>,
<matplotlib.axis.XTick at 0x1fda0046e10>,
<matplotlib.axis.XTick at 0x1fd9ffab278>,
<matplotlib.axis.XTick at 0x1fda007ee80>,
<matplotlib.axis.XTick at 0x1fda008e4a8>,
<matplotlib.axis.XTick at 0x1fda008e9b0>],
<a list of 6 Text xticklabel objects>)
```

```
Out[185]: []
```

```
Out[185]: []
```

```
Out[185]: <matplotlib.legend.Legend at 0x1fda008e320>
```



SVR with Kernel as 'rbf' Result:

Best parameters: {'C': 100, 'gamma': 0.1}

Average Cross validation score: -0.0372

Test score: -0.0345

Best Model for the prediction

Based on the cross validation score and the test score for above models, it is inferred that the polynomial regression is the best model to predict the house prices.

```
In [201]: d = {'Model': ['Linear Regression', 'KNN Regression', 'Ridge Regression', 'Lasso Regression', 'Polynomial Regression', 'Simple SVR', 'SVR with Linear kernel', 'SVR with Poly kernel', 'SVR with rbf kernel'],
              'Cross-Validation Score': [grid_search_lr.best_score_, grid_search_knn.best_score_, grid_search_rr.best_score_, grid_search_lasso.best_score_, grid_search_poly.best_score_, grid_search_svr.best_score_, grid_search_svr_linear.best_score_, grid_search_svr_poly.best_score_, grid_search_svr_rbf.best_score_]}
```

```
In [204]: result = pd.DataFrame(data=d)
result
```

```
Out[204]:
```

	Model	Cross-Validation Score
0	Linear Regression	0.703833
1	KNN Regression	0.705060
2	Ridge Regression	0.704406
3	Lasso Regression	0.704398
4	Polynomial Regression	0.791462
5	Simple SVR	0.554950
6	SVR with Linear kernel	0.631385
7	SVR with Poly kernel	0.206212
8	SVR with rbf kernel	-0.041249

```
In [187]: pol = PolynomialFeatures(degree = 2)
X_pol = pol.fit_transform(X_train_org)
Xt_pol = pol.fit_transform(X_test_org)
pol_reg = LinearRegression()
pol_reg.fit(X_pol, y_train)
ypred = pol_reg.predict(Xt_pol)
```

```
Out[187]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```



```
In [195]: with np.printoptions(threshold=np.inf):
           print(ypred)
```

```
[ 341597.40040728  702660.62726908  484603.9052442   429292.390924
  378652.15614089  291040.97503573  910605.41031334  921098.78258406
  460615.36765849  586550.24702009  731334.69904633  394656.6315059
  475202.15136696  426664.17566913  822988.67983069  758345.93563601
  654118.67215681  666314.05665389  298052.69265396  590541.03637684
  559036.23678609  385056.74812627  221439.20806613  547969.07705364
  337678.51230448  753764.27594396  429842.28582262  224481.99298393
  478859.63070237  431387.78808331  489926.17724677  414933.22694872
  292005.74726543  332530.19191765  852038.64820287  265889.67016193
  316979.95360975  262764.64894809  548419.99999253  384401.00567746
  457103.39279147  750351.72843836  1187277.64062186  794662.97705005
  426535.9462004   397197.19972121   689789.19742275  474511.31364627
  379114.87023796  883924.68498636  1041520.18617112  751244.89115981
  513184.7653153   339289.05826937   333057.43847961  479922.42780261
  436867.24486438  231423.59747122   268730.40000508  720082.15212843
  448312.81380294  651346.79590288   554844.33883619  762787.29889828
  962364.68571667  534996.54897123   345524.12955321  600983.06320147
  257289.54212081  288191.6280486    921666.82460959  233585.0619591
 1260677.75732788  338789.54883157   274001.08604076  540877.58997174
  330702.53700000  330702.53700000  330702.53700000  330702.53700000]
```